

Datacenter Technology (Fall, 2018)

Software-Defined Data Center

Lingfang Zeng (曾令仿)

Wuhan National Laboratory for Optoelectronics (WNLO)

Huazhong University of Science and Technology (HUST)

<https://lingfangzeng.github.io/>



Outline



SDN

**SOFTWARE DEFINED
NETWORKS**



UOS

**UBIQUITOUS
OPERATING SYSTEMS**



SDFS

**SOFTWARE-DEFINED
FILE SYSTEM**

Thanks

- Scott Shenker. *The Future of Networking, and the Past of Protocols*. Fall 2011.
- Hong Mei, Yao Guo. *Toward Ubiquitous Operating Systems: A Software-Defined Perspective*. IEEE Computer 51(1): 50-56 (2018)
- Jiahao Liu, Fang Wang, Lingfang Zeng, Dan Feng, Tingwei Zhu. *SDFS: A Software-Defined File System for Multi-Tenant Cloud Storage*. Software: Practice and Experience. (Accepted, 2018) (CCF B)



01

SDN

SOFTWARE DEFINED NETWORKS

Key to Internet Success: Layers

Applications

...built on...

Reliable (or unreliable) transport

...built on...


Best-effort global packet delivery

...built on...

Best-effort local packet delivery

...built on...

Physical transfer of bits



Why Is Layering So Important?

- Decomposed delivery into fundamental components
- Independent but compatible innovation at each layer
- A practical success of unprecedented proportions...
- ...but an academic failure

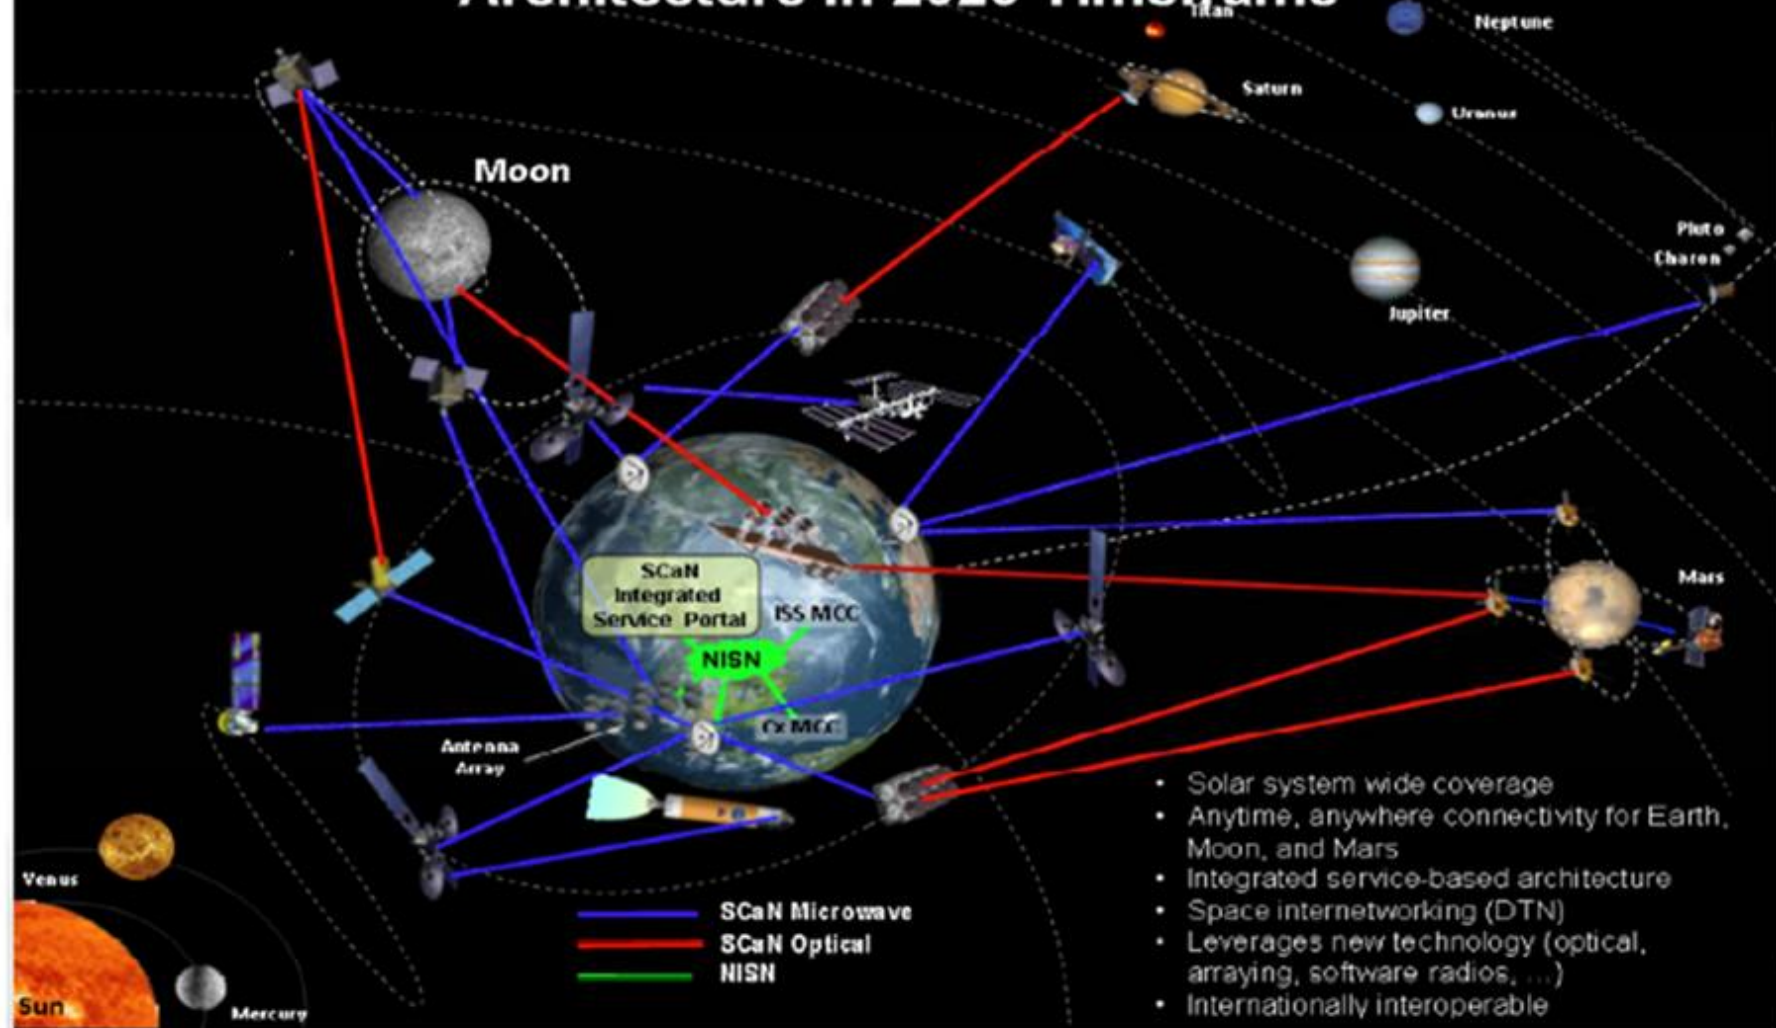
Built an Artifact, Not a Discipline

- Other fields in “systems”: OS, DB, DS, etc.
 - Teach basic principles
 - Are easily managed
 - Continue to evolve
- Networking
 - Teach big bag of protocols
 - Notoriously difficult to manage
 - Evolves very slowly

Why Does Networking Lag Behind?

- Networks used to be simple: Ethernet, IP, TCP....
- New **control** requirements led to great complexity
 - Isolation → VLANs, ACLs
 - Traffic engineering → MPLS, ECMP, Weights
 - Packet processing → Firewalls, NATs, middleboxes
 - Analysis → Deep packet inspection (DPI)
- Mechanisms designed and deployed independently
 - Complicated “control plane” design, primitive functionality
 - Stark contrast to the elegantly modular “data plane”

SCaN Notional Integrated Communication Architecture in 2025 Timeframe



Layers are Great Abstractions

- Layers only deal with the data plane
- We have no powerful control plane abstractions!
- How do we find those control plane abstractions?
- Two steps: define problem, and then **decompose** it.

The Network Control Problem

- Compute the configuration of each physical device
 - E.g., Forwarding tables, ACLs,...
- Operate without communication guarantees
- Operate within given network-level protocol

Only people who love complexity would find this a reasonable request

From Requirements to Abstractions

- Operate without communication guarantees
 - Need an abstraction for distributed state
- Compute the configuration of each physical device
 - Need an abstraction that simplifies configuration
- Operate within given network-level protocol
 - Need an abstraction for general forwarding model

Once these abstractions are in place, control mechanism has a much easier job!

(1) Distributed State Abstraction


- Shield control mechanisms from state distribution
 - While allowing access to this state
- Natural abstraction: global network view
 - Annotated network graph provided through an API
- Implemented with “Network Operating System”
 - Control mechanism is now program using API
- No longer a distributed protocol, now just a graph algorithm
 - E.g. Use Dijkstra rather than Bellman-Ford

Network of Software Defined Network Routers

e.g. routing, access control


Control Program

Global Network View



Distributed algorithm running between neighbors

Network OS



Major Change in Paradigm

- No longer designing distributed control protocols
- Design one distributed system (NOS)
- Use for all control functions
- Now just defining a centralized control function
- Configuration = Function(view)



(2) Specification Abstraction

- Control program should express desired behavior
- It should not be responsible for implementing that behavior on physical network infrastructure
- Natural abstraction: **simplified model of network**
 - Simple model with only enough detail to specify goals
- Requires a new shared control layer
 - Map abstract configuration to physical configuration
- This is “network virtualization”

Simple Example: Access Control

What

Abstract
Network
Model



Global
Network
View


How

Software Defined Network: Take 2

Abstract Network Model




Global Network View



What Does This Picture Mean?

- Write a simple program to configure a simple model
 - Configuration merely a way to specify what you want
- Examples
 - ACLs: who can talk to who
 - Isolation: who can hear my broadcasts
 - Routing: only specify routing to the degree you care
 - Some flows over satellite, others over landline
 - TE: specify in terms of quality of service, not routes
- Virtualization layer “compiles” these requirements
 - Produces suitable configuration of actual network devices
- NOS then transmits these settings to physical boxes

Software Defined Network: Take 2



SDN: Clean Separation of Concerns

- Control prgm: specify behavior on abstract model
 - Driven by Operator Requirements
- Net Virt'n: map abstract model to global view
 - Driven by **Specification** Abstraction
- NOS: map global view to physical switches
 - API: driven by **Distributed State** Abstraction
 - Switch/fabric interface: driven by **Forwarding** Abstraction

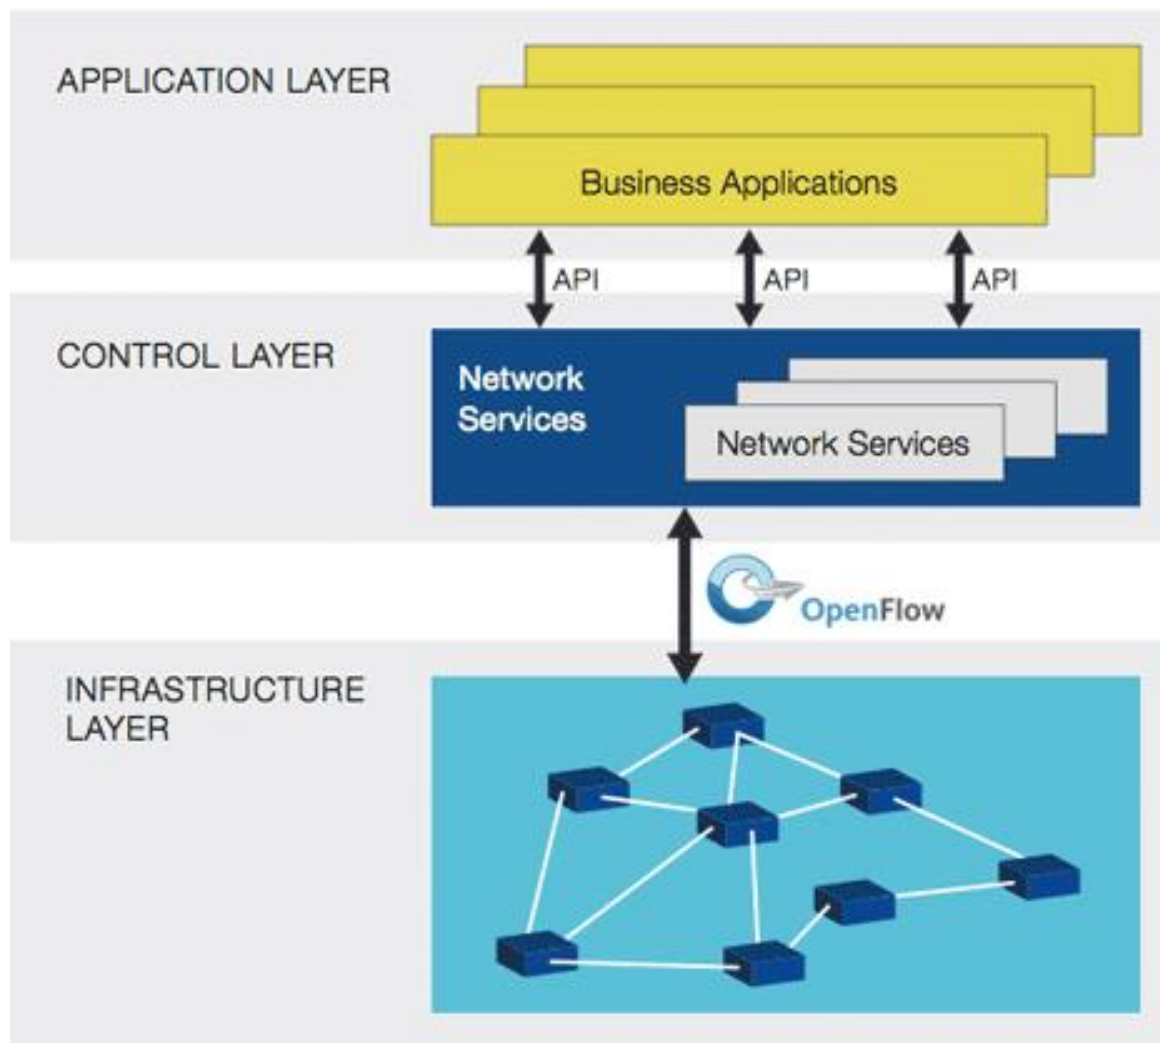
(2) Forwarding Abstraction

- Switches have two “brains”
 - Management CPU (smart but slow)
 - Forwarding ASIC (fast but dumb)
- Need a forwarding abstraction for both
 - CPU abstraction can be almost anything
- ASIC abstraction is much more subtle: OpenFlow
- OpenFlow
 - Control switch by inserting <header;action> entries
 - Essentially gives NOS remote access to forwarding table
 - Instantiated in OpenvSwitch

SDN: Clean Separation of Concerns

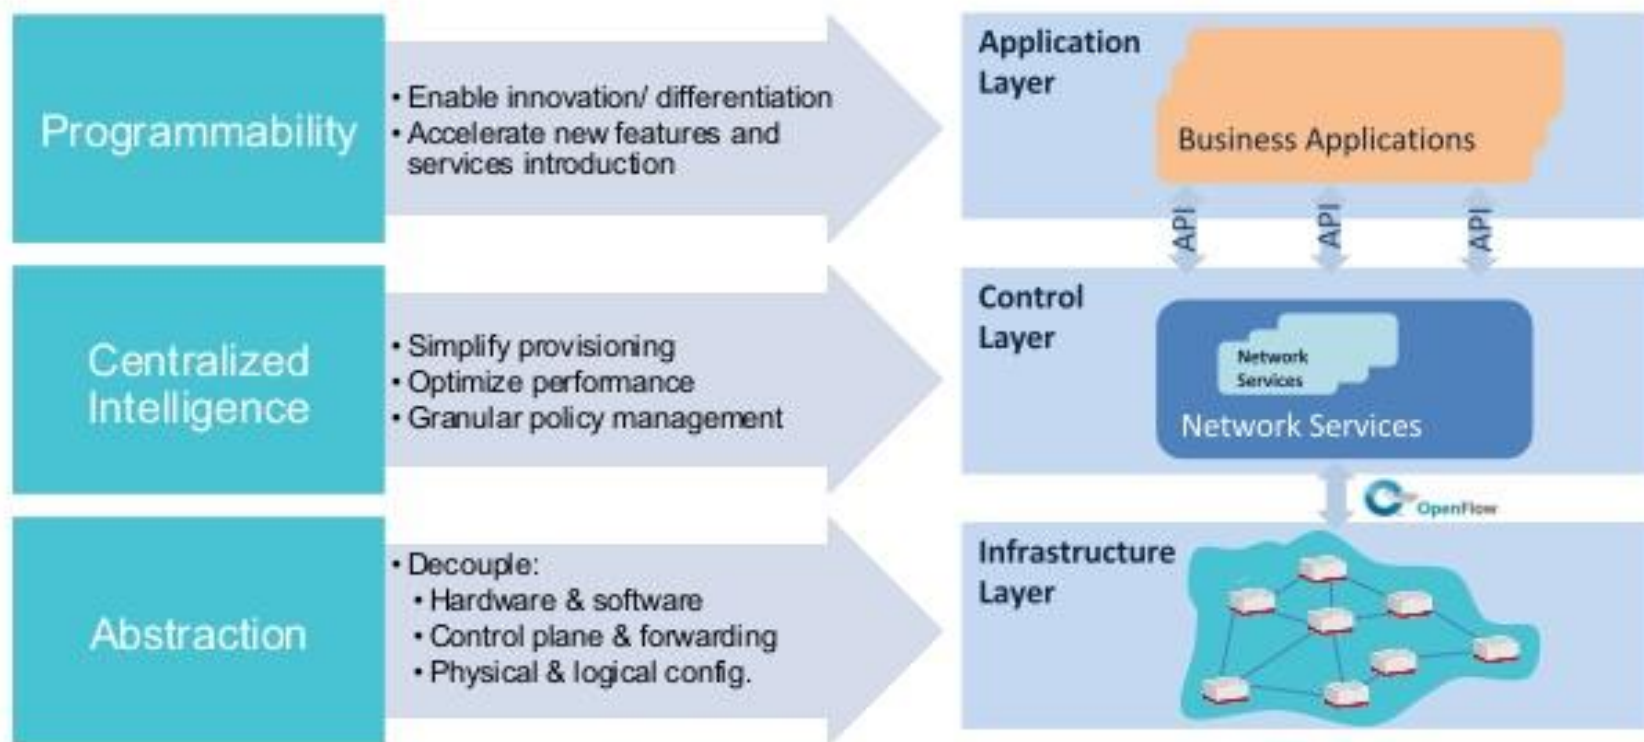
- Control prgm: specify behavior on abstract model
 - Driven by Operator Requirements
- Net Virt'n: map abstract model to global view
 - Driven by **Specification** Abstraction
- NOS: map global view to physical switches
 - API: driven by **Distributed State** Abstraction
 - Switch/fabric interface: driven by **Forwarding Abstraction**

SDN - Open Networking Foundation

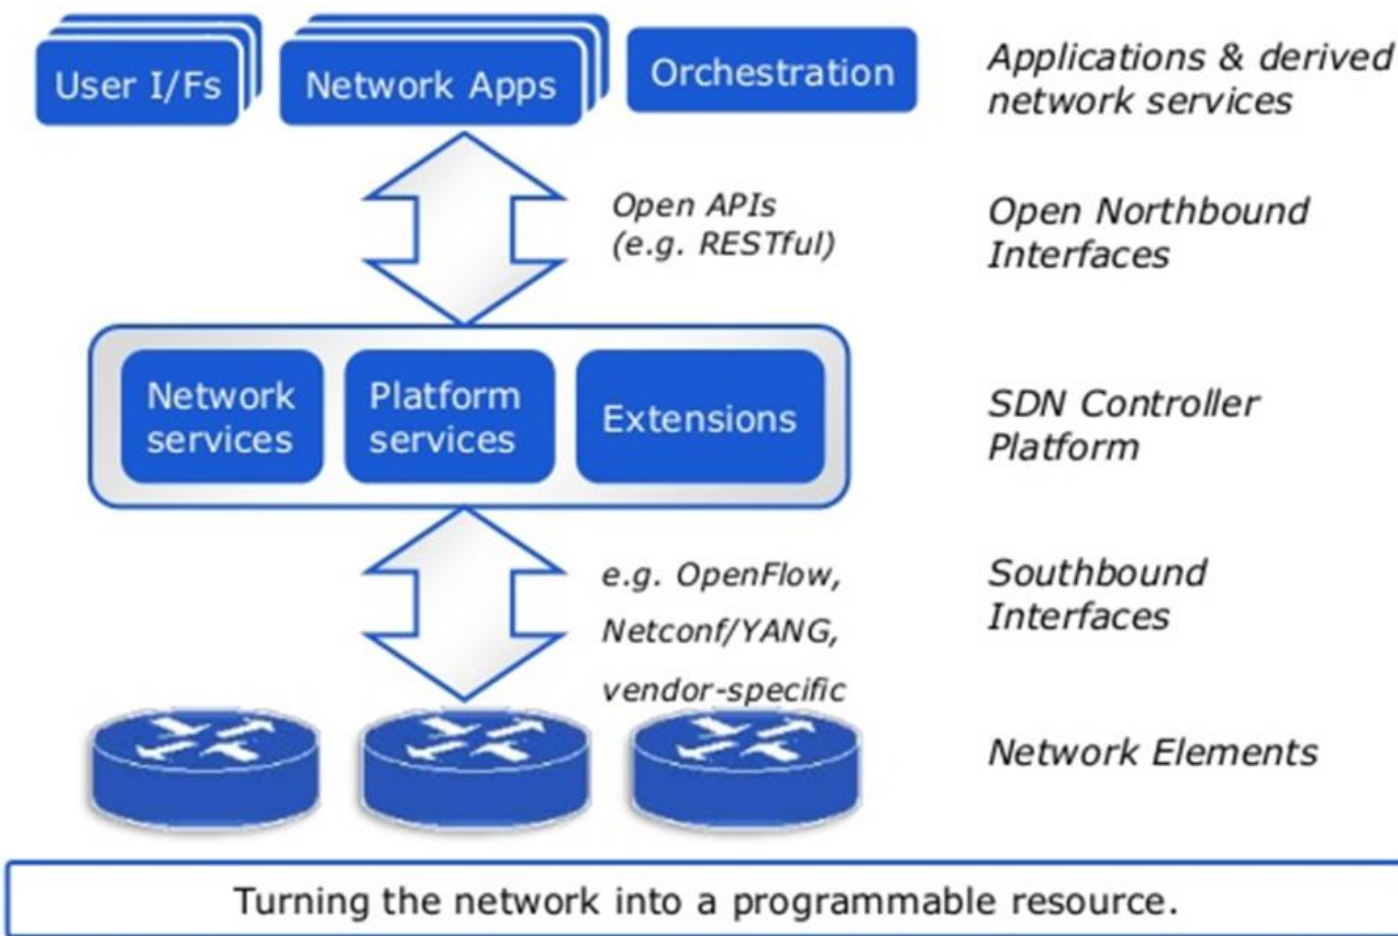


SDN - Open Networking Foundation

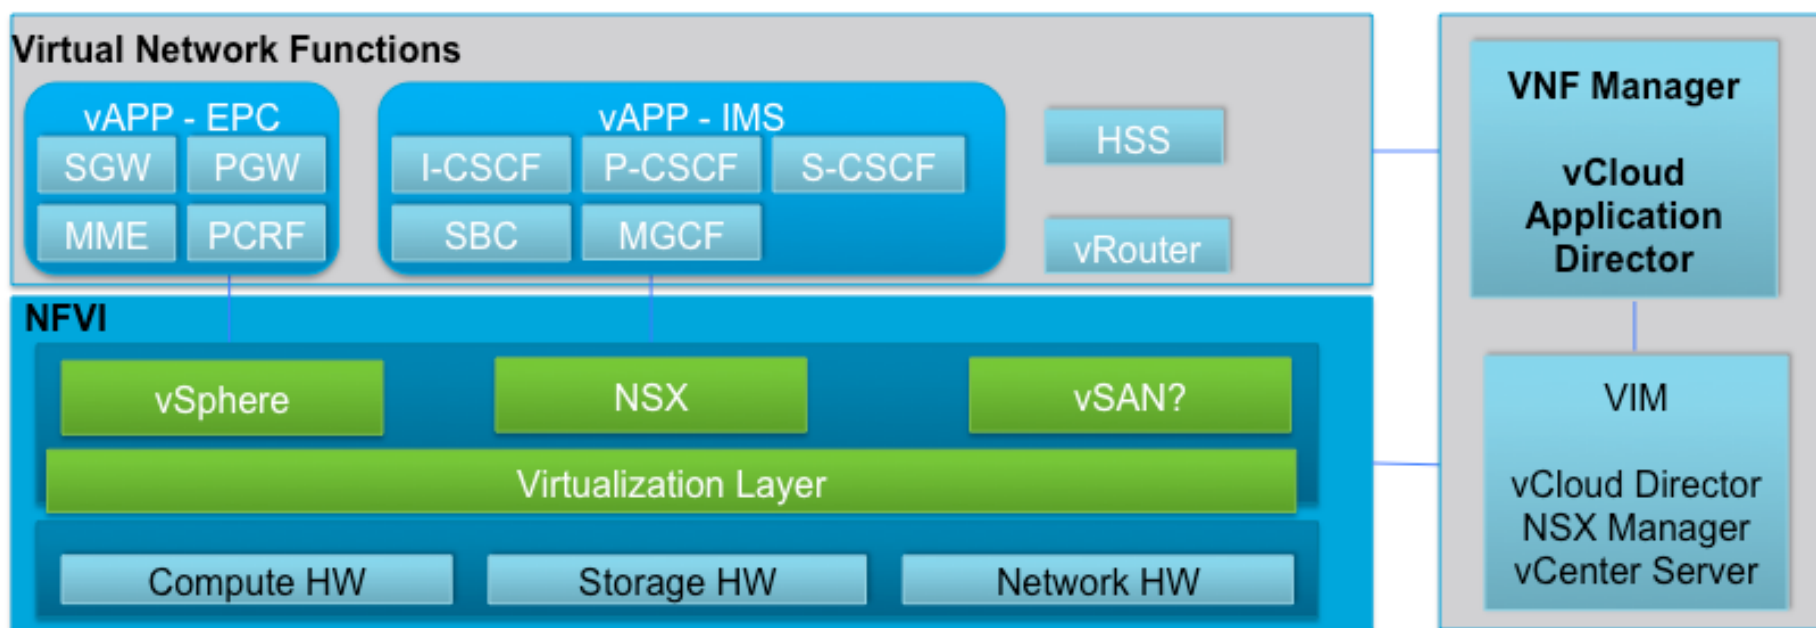
ONF SDN Architecture



SDN - Open Networking Foundation

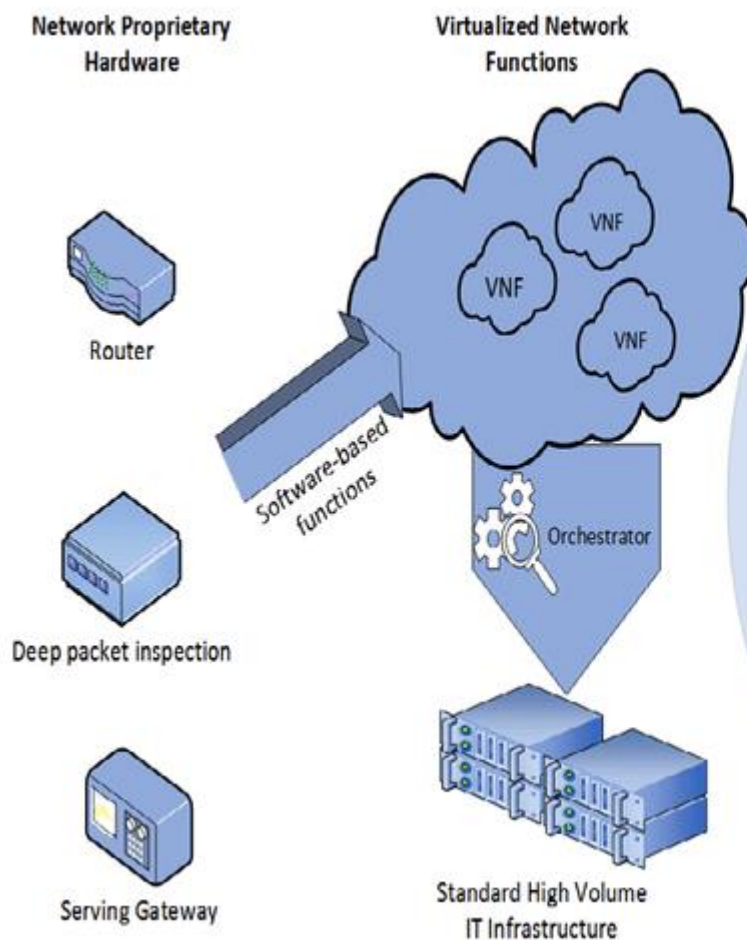


SDN vs. Virtual Network Functions



SDN vs. Virtual Network Functions

Software-Defined Networking (SDN) and NFV are two different independent technologies; however, they are complementary to each other.



NFV

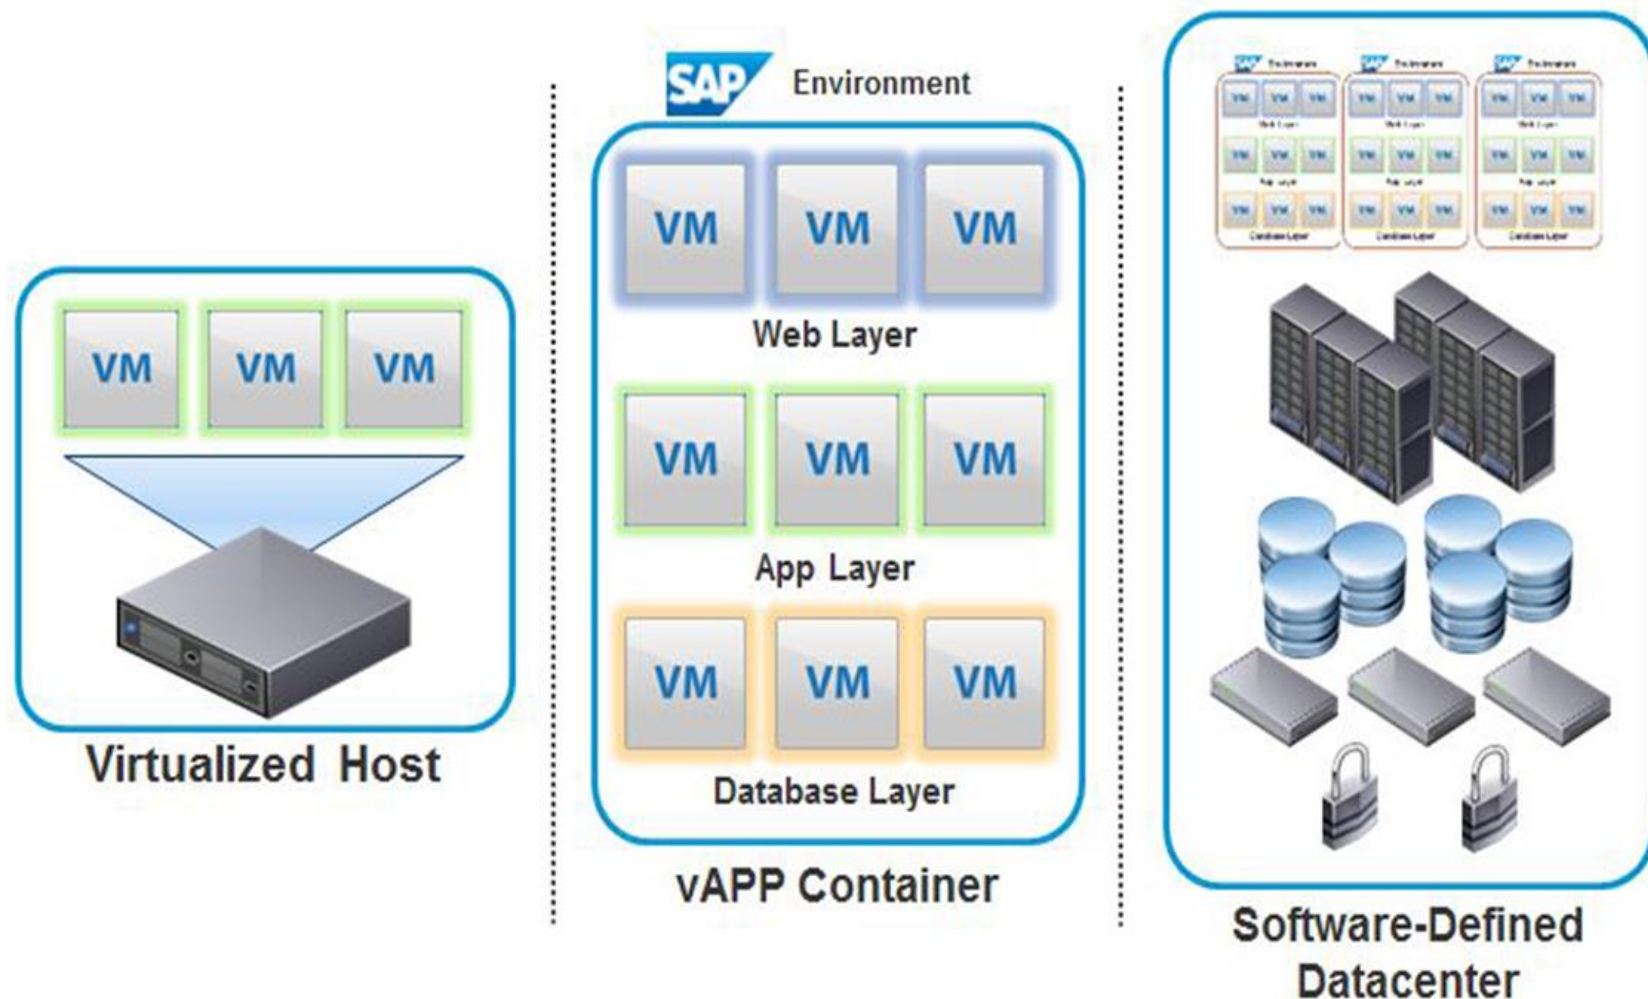
SDN

1) NFV is the concept of transferring the network functions from dedicated hardware appliances to software-based applications.
2) NFV decouples the network functions from the proprietary hardware without affecting the functionality.

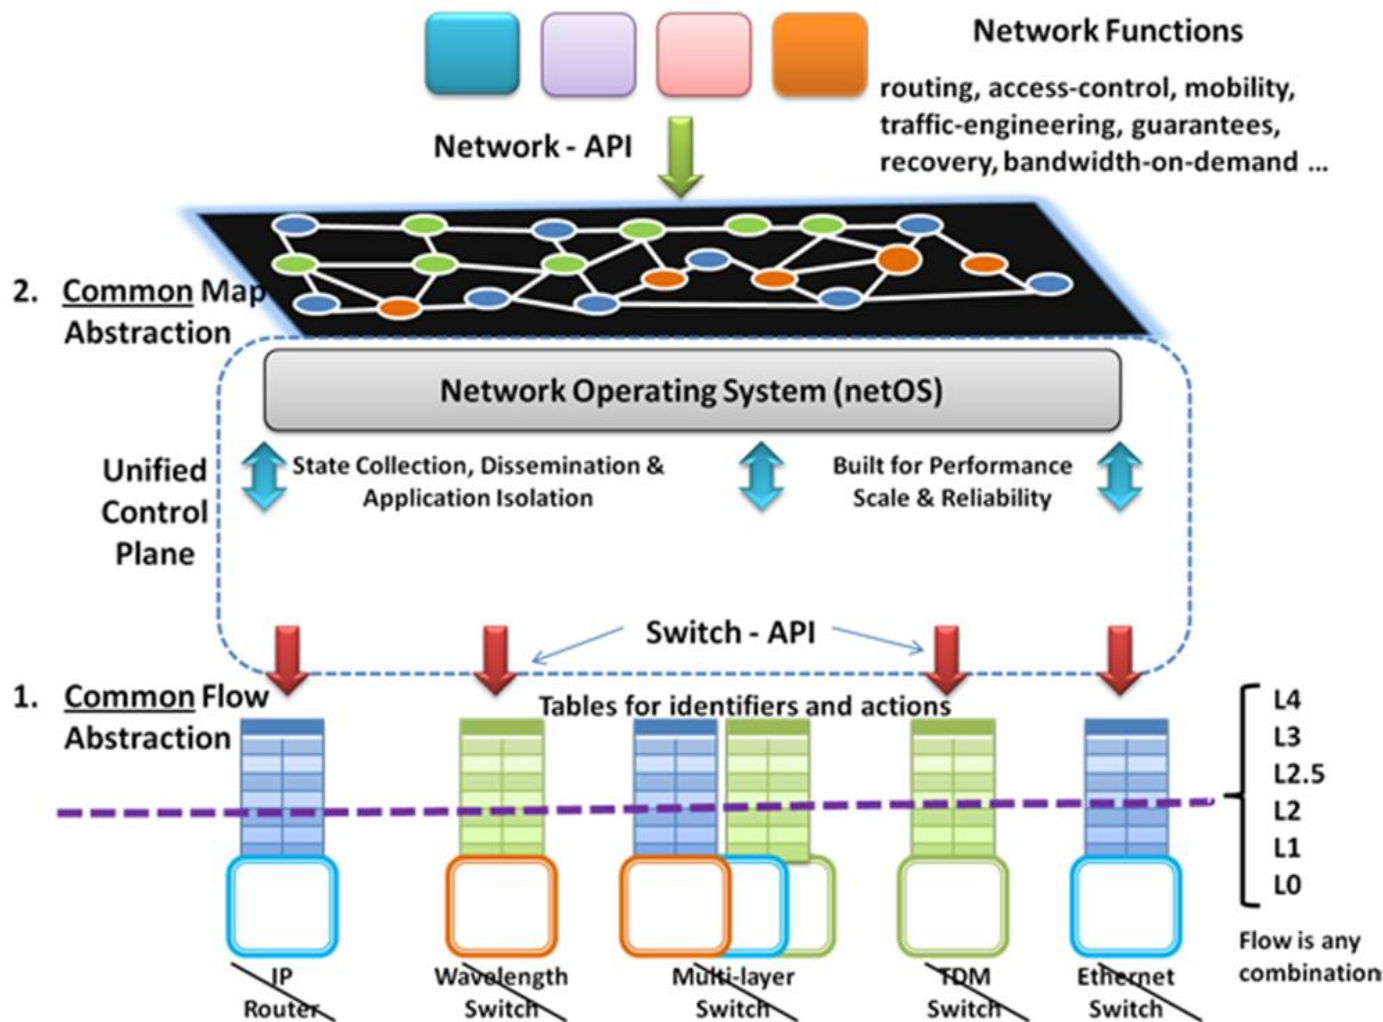
1) SDN serve NFV by providing the programmable connectivity between VNFs; these connections can be managed by the orchestrator of the VNFs which will mimic the role of the SDN controller [6].
2) NFV serve SDN by implementing its network functions in a software manner on a COTS servers. It can virtualize the SDN Controller to run on cloud, which could be migrated to best fit location according to the network needs.

1) SDN is concerned in network functionalities it decouples the control and data planes[6].
2) SDN Provides centralized controller and network programmability[6].

Software-Defined-Anything (SDx)



Software-Defined-Anything (SDx)

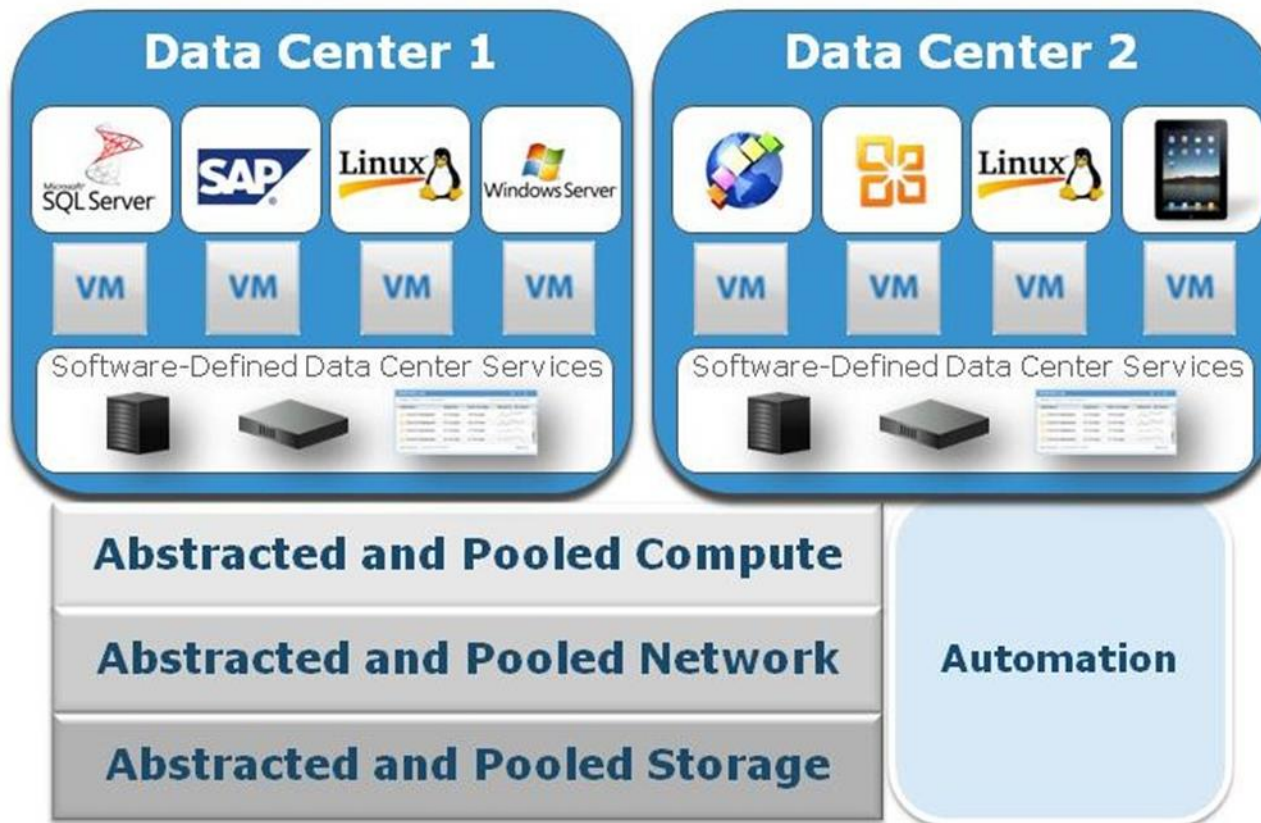


Software-Defined-Anything (SDx)




Software-Defined-Anything (SDx)

The Software-Defined Data Center



Software-Defined-Anything (SDx)





02

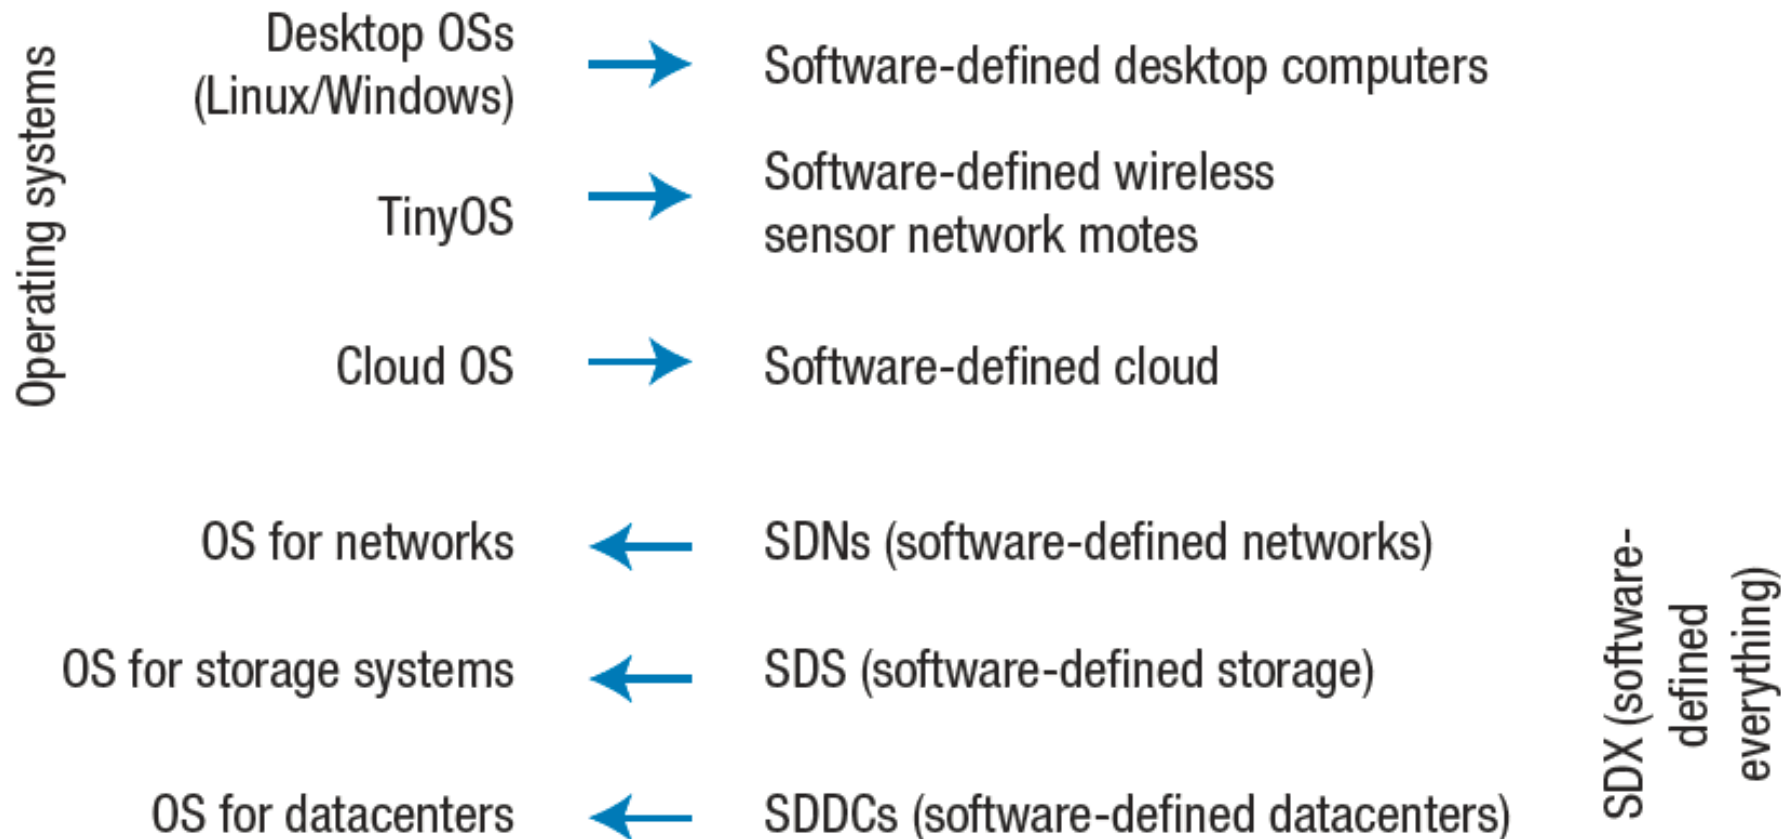
UOS

UBIQUITOUS OPERATING SYSTEMS

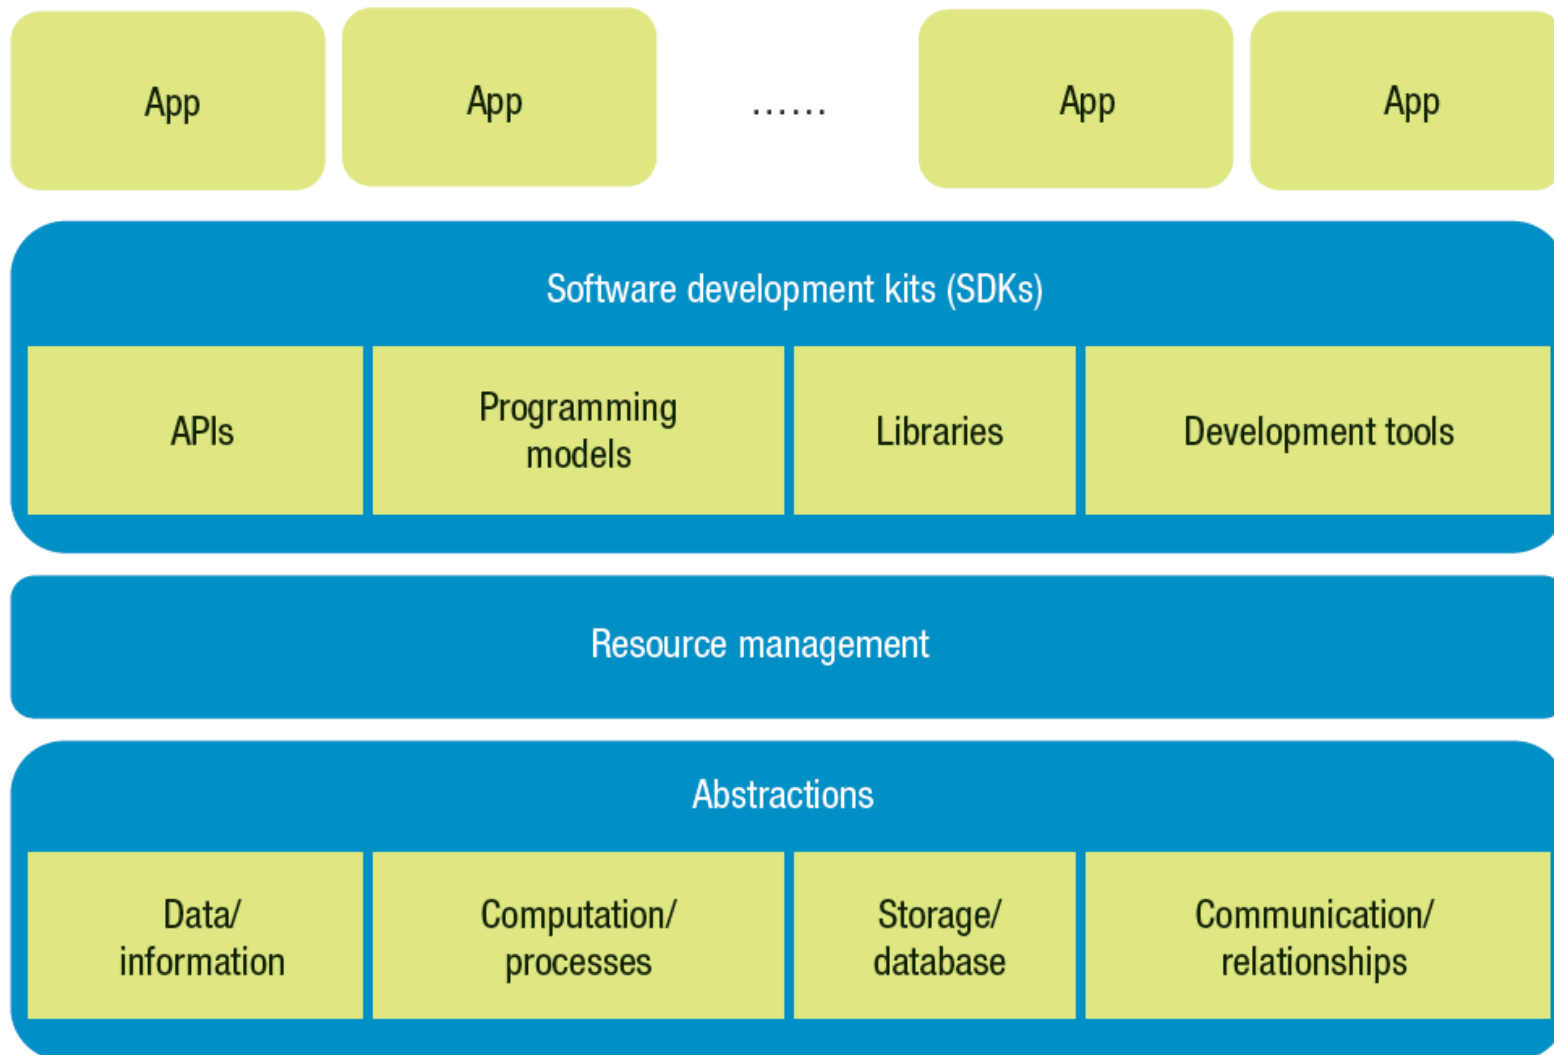
Case Study 1 - Toward Ubiquitous Operating Systems

Timeframe	Representative OS(s)	Computer system	Main characteristics
1956	GM-NAA I/O	IBM 704	The first practical OS Simple batch processing I/O management
1960s	IBM OS/360 series	IBM 360 series—mainframes	Time-sharing Multibatch processing Memory management Virtual machines (VM/370)
1970s	Unix	Minicomputers/workstations	First modern OS Developed with machine-independent languages (C) Provides standard interfaces Integrated development environment
1980s	Mac OS, Windows, Linux	Personal computers (PCs)	Provides modern GUI Improves usability for personal users
2000s	Apple iOS, Google Android, Windows Phone	Smartphones	Customization of traditional OSs Improves usability for mobile devices New app delivery model (App Store, Google Play)

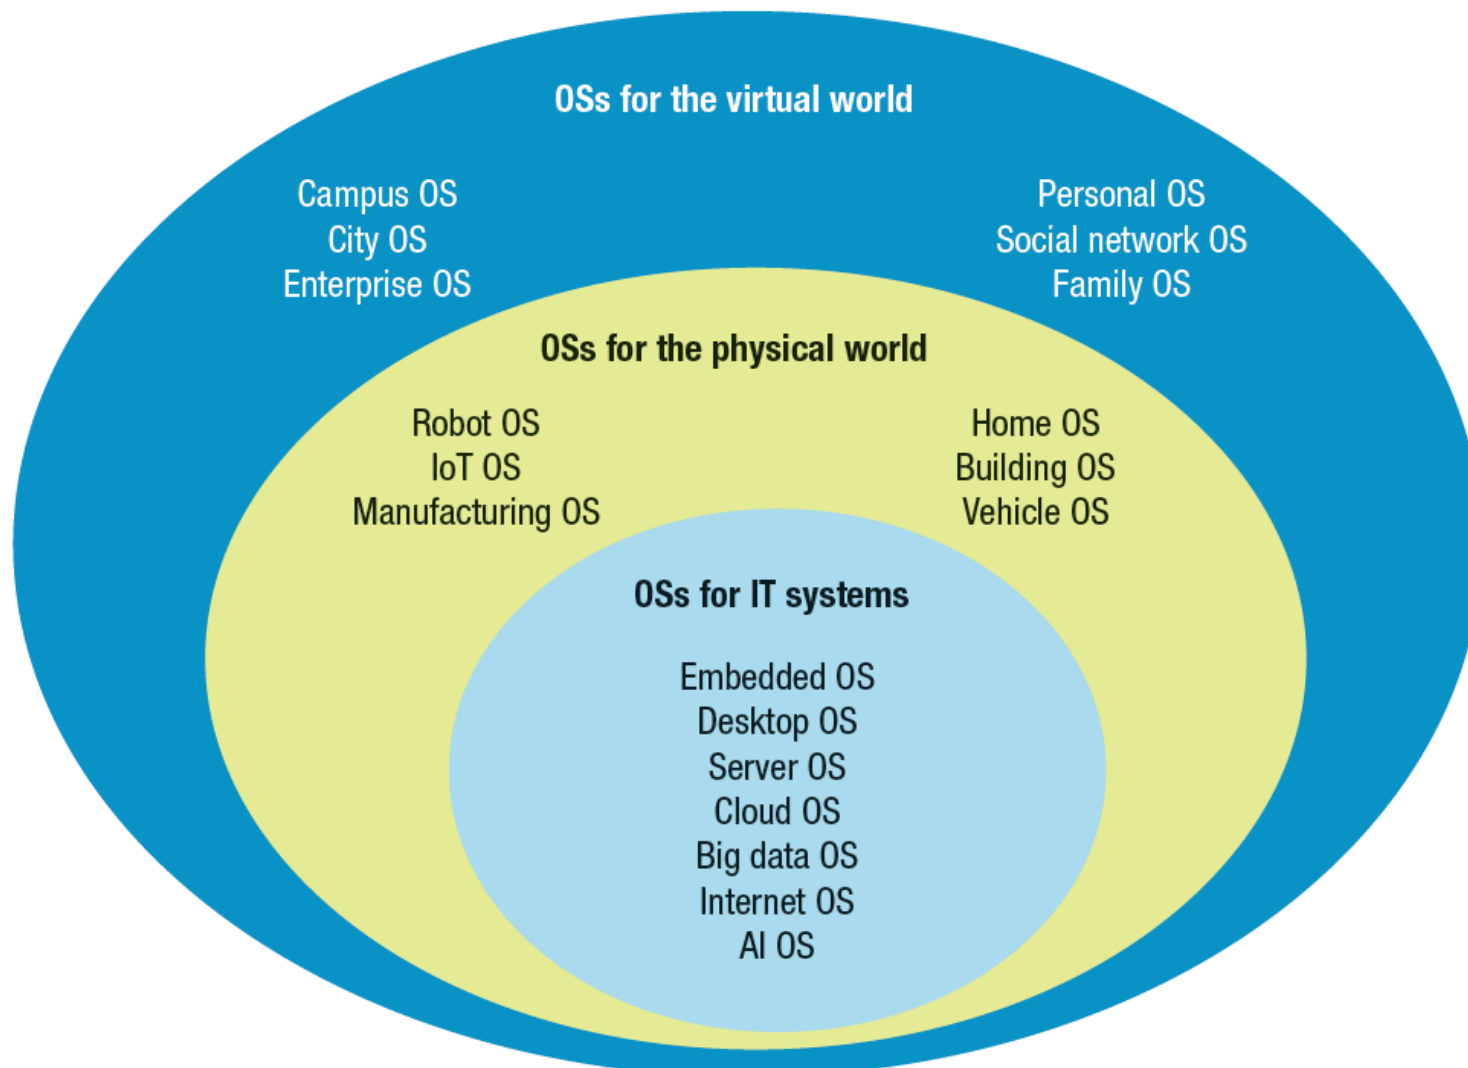
Toward Ubiquitous Operating Systems



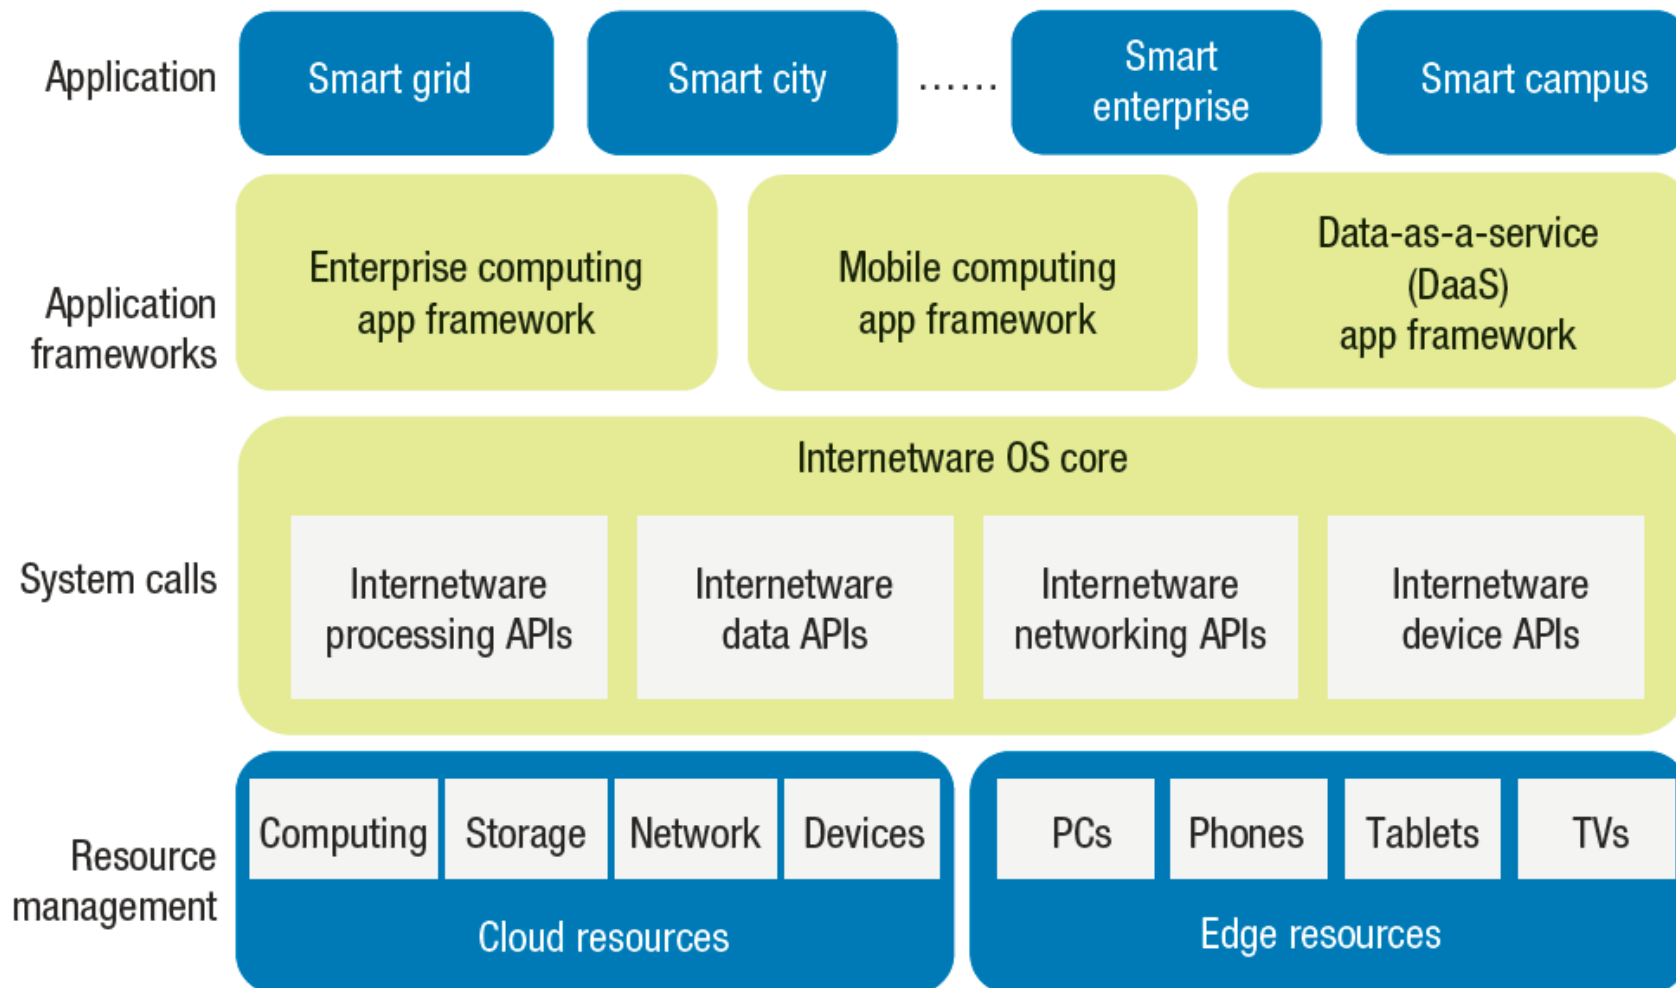
Toward Ubiquitous Operating Systems



Toward Ubiquitous Operating Systems



Toward Ubiquitous Operating Systems



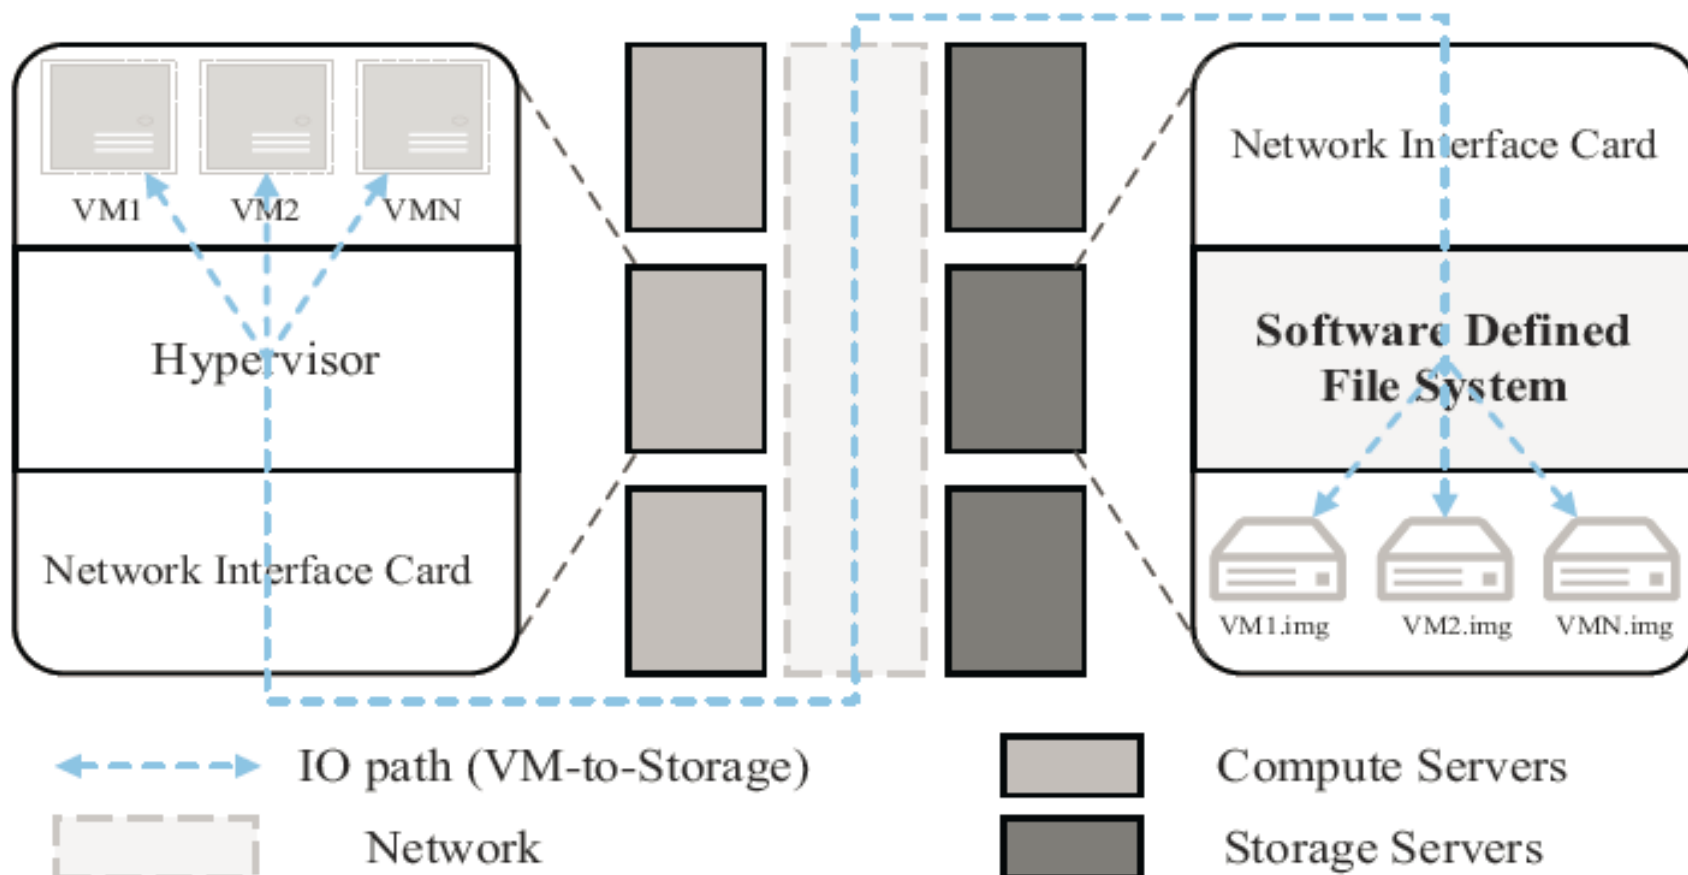


03

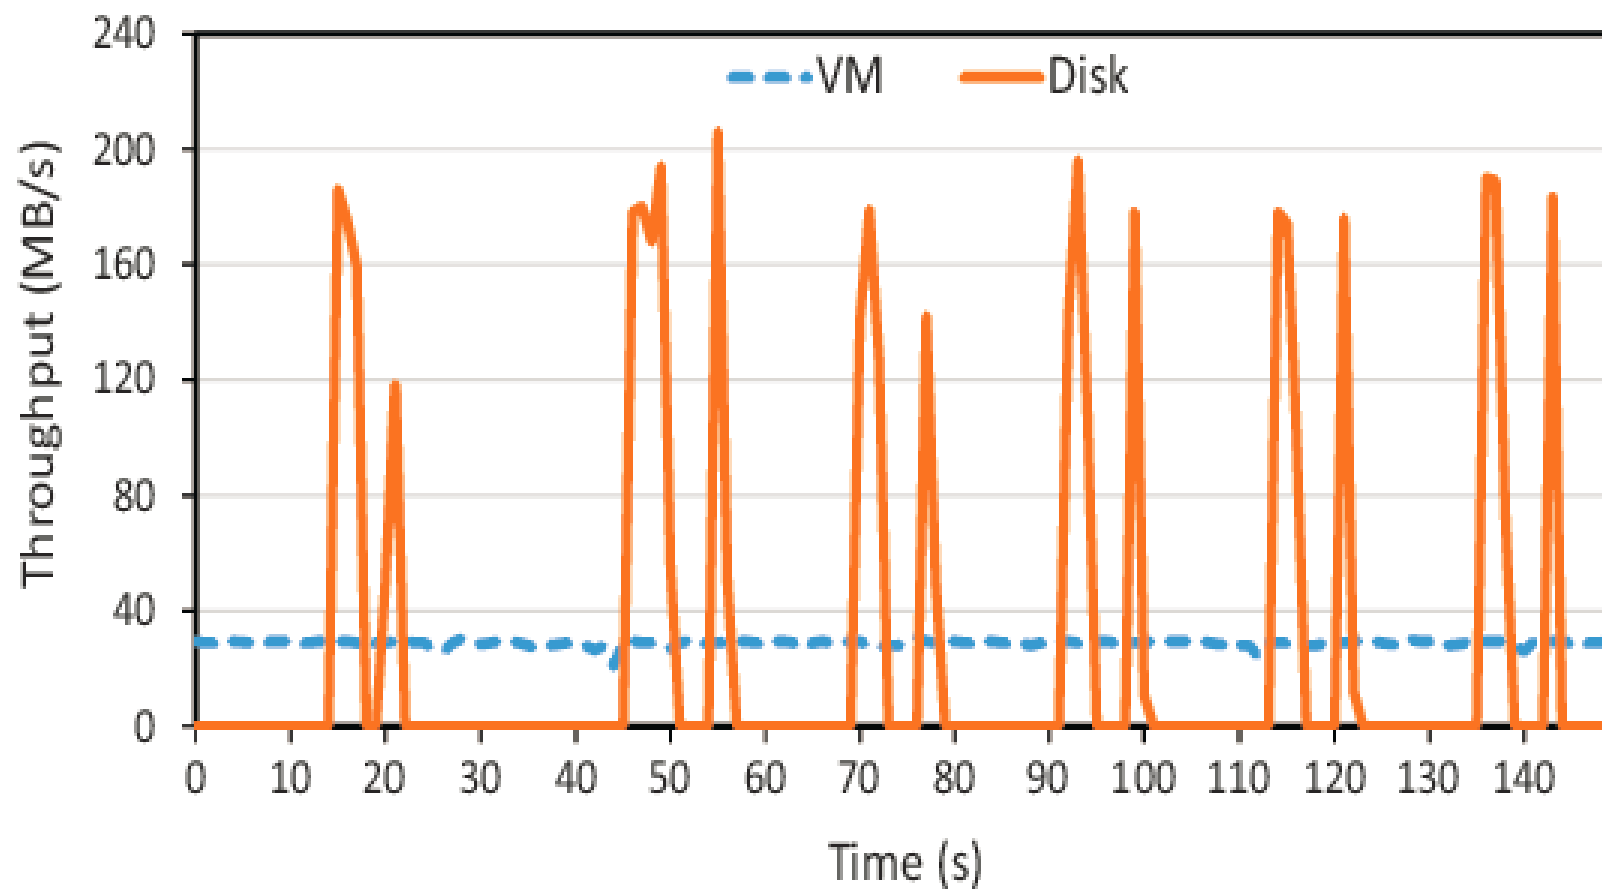
SDFS

SOFTWARE-DEFINED FILE SYSTEM

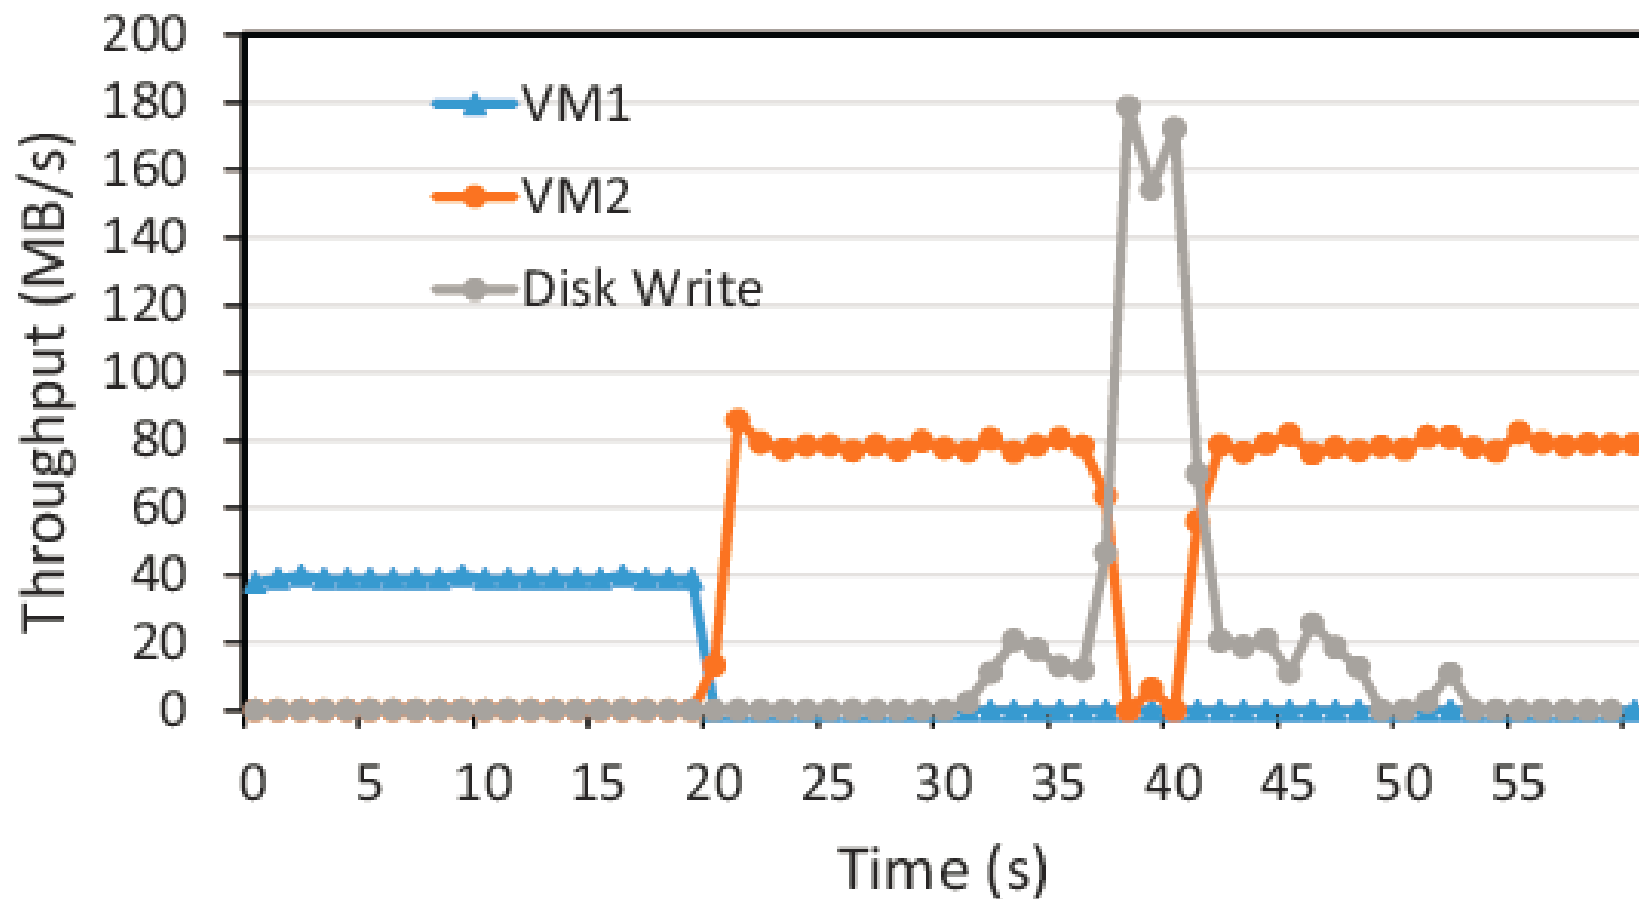
Case Study 2 - Software-Defined File System




Motivation




Motivation




Motivation



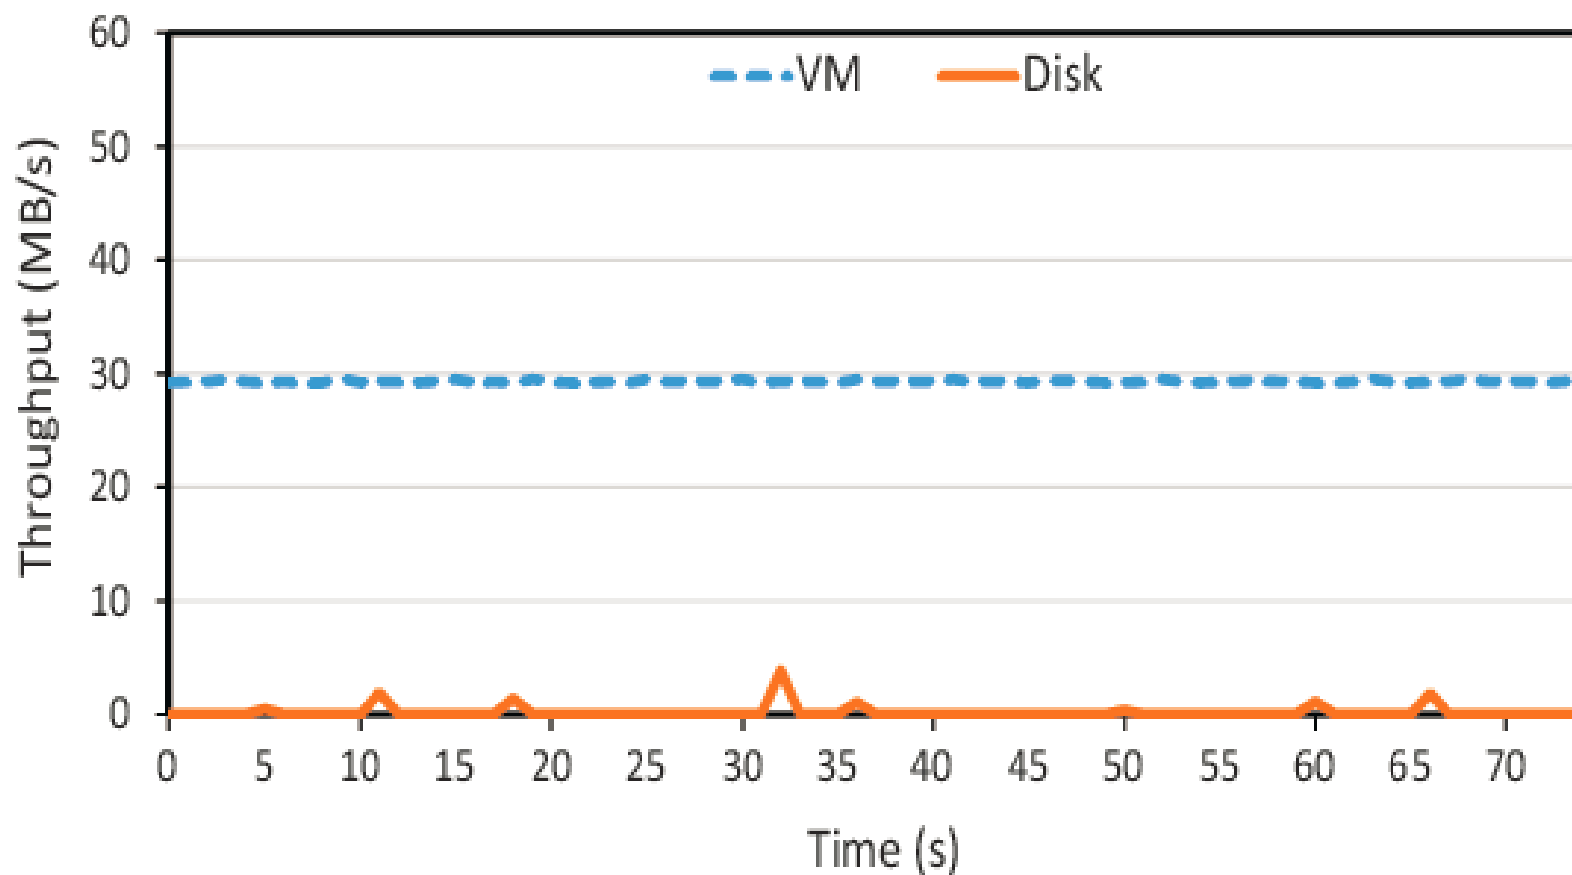
Motivation




Motivation



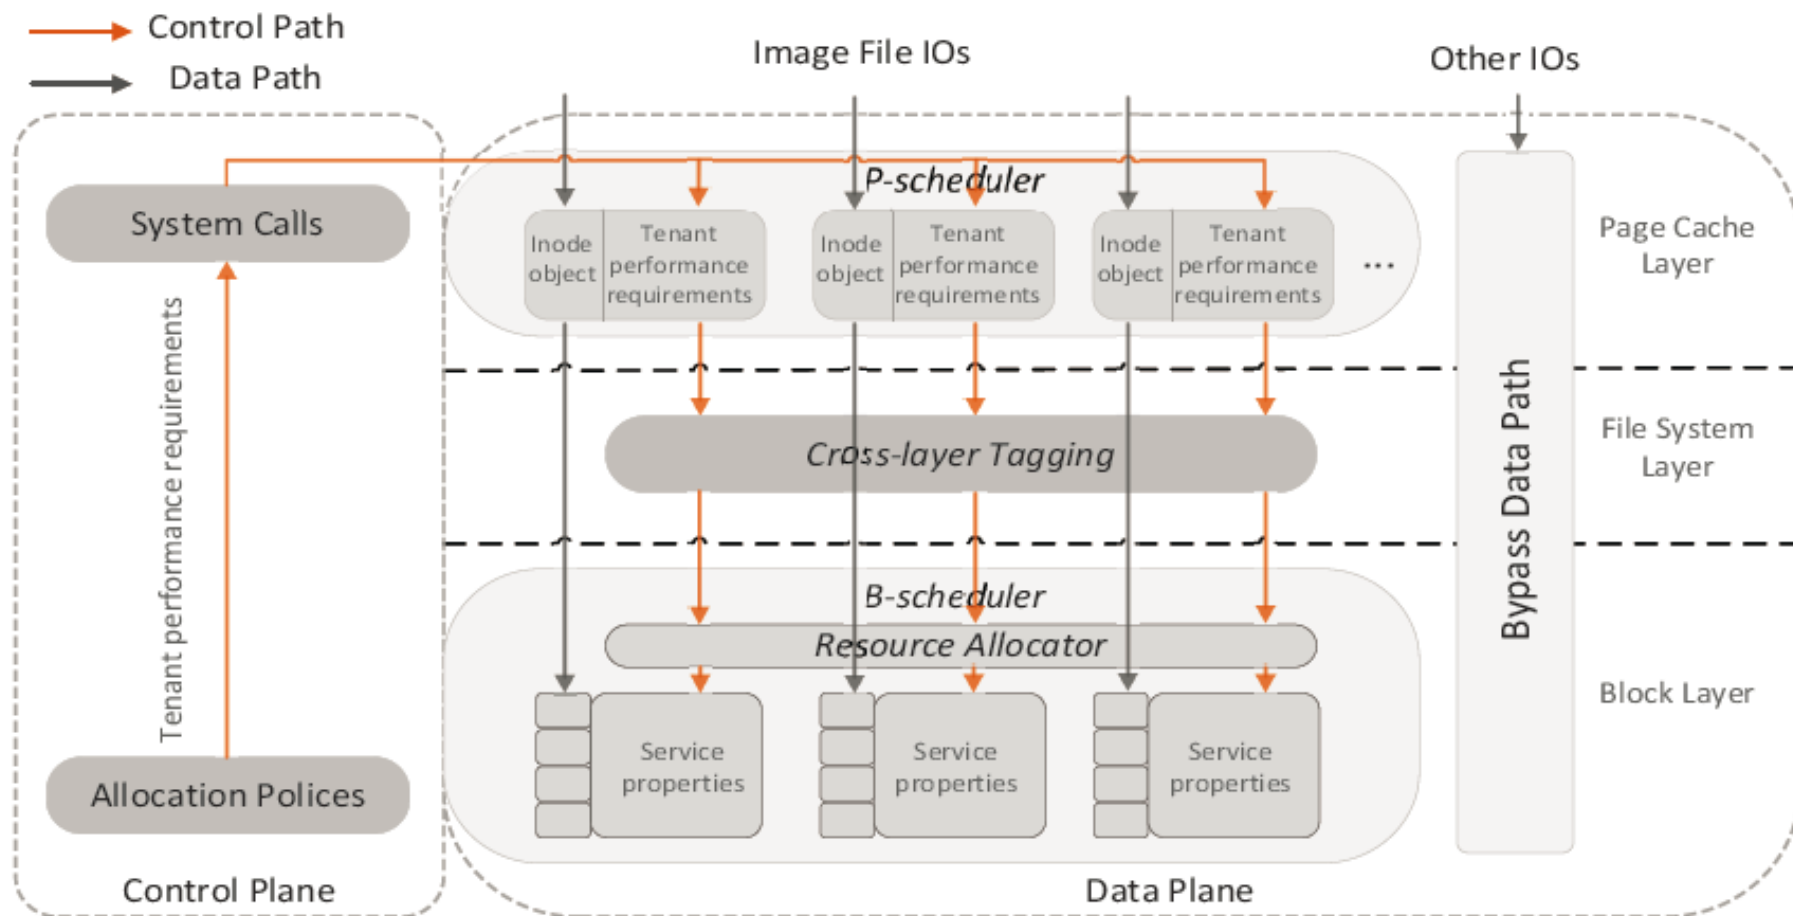
Motivation



Motivation



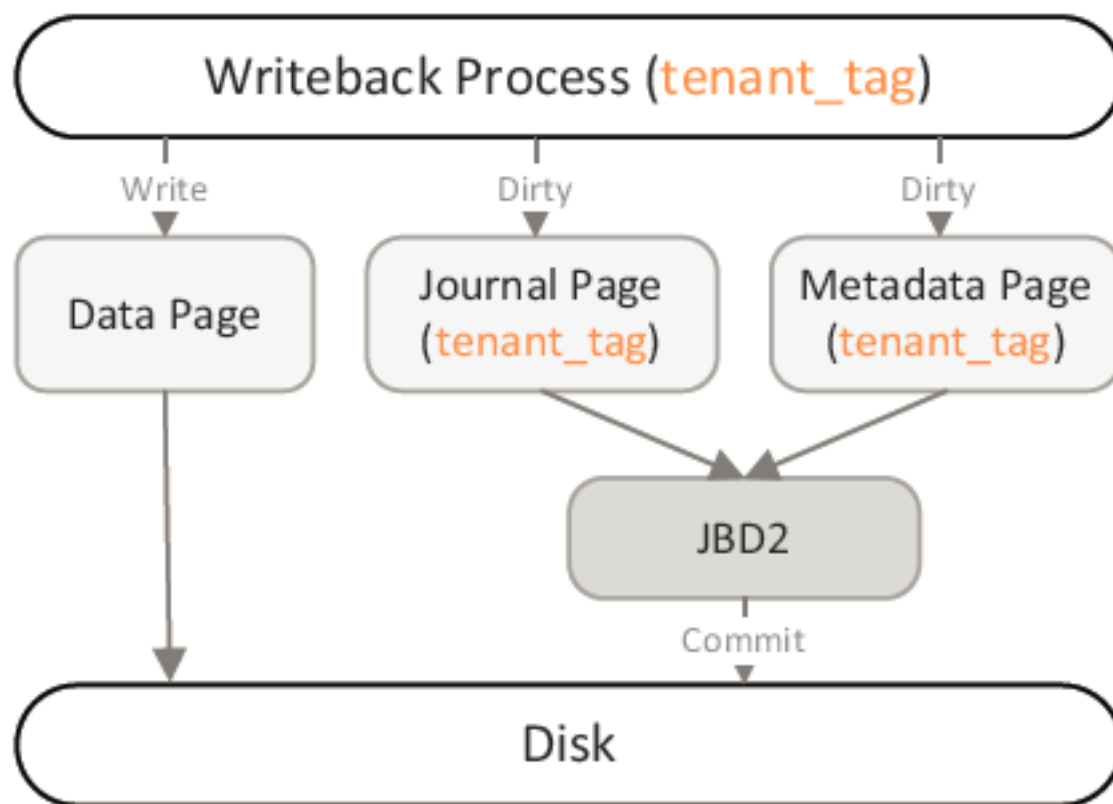
SDFS: Design



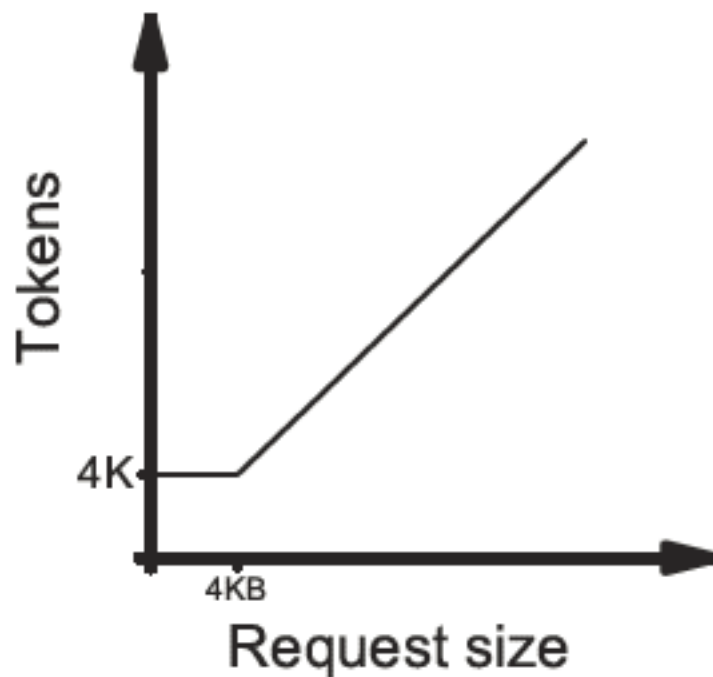
SDFS: Control Plane

Function	Description
<i>set_tenant_tag_fixed</i>	sets a fixed throughput guarantee (f) for a VM. $f \in (0, 200)$
<i>set_tenant_tag_min</i>	sets a minimum throughput guarantee (m) for a VM. $m \in (0, 200)$
<i>set_tenant_tag_max</i>	sets an upper limit on resource usage (l) for a VM. $l \in (0, 200)$
<i>set_tenant_tag_infi</i>	does not set resource usage limit for a VM.
<i>get_tenant_tag</i>	returns tenant performance requirements


SDFS: Control Plane



SDFS: Data Plane

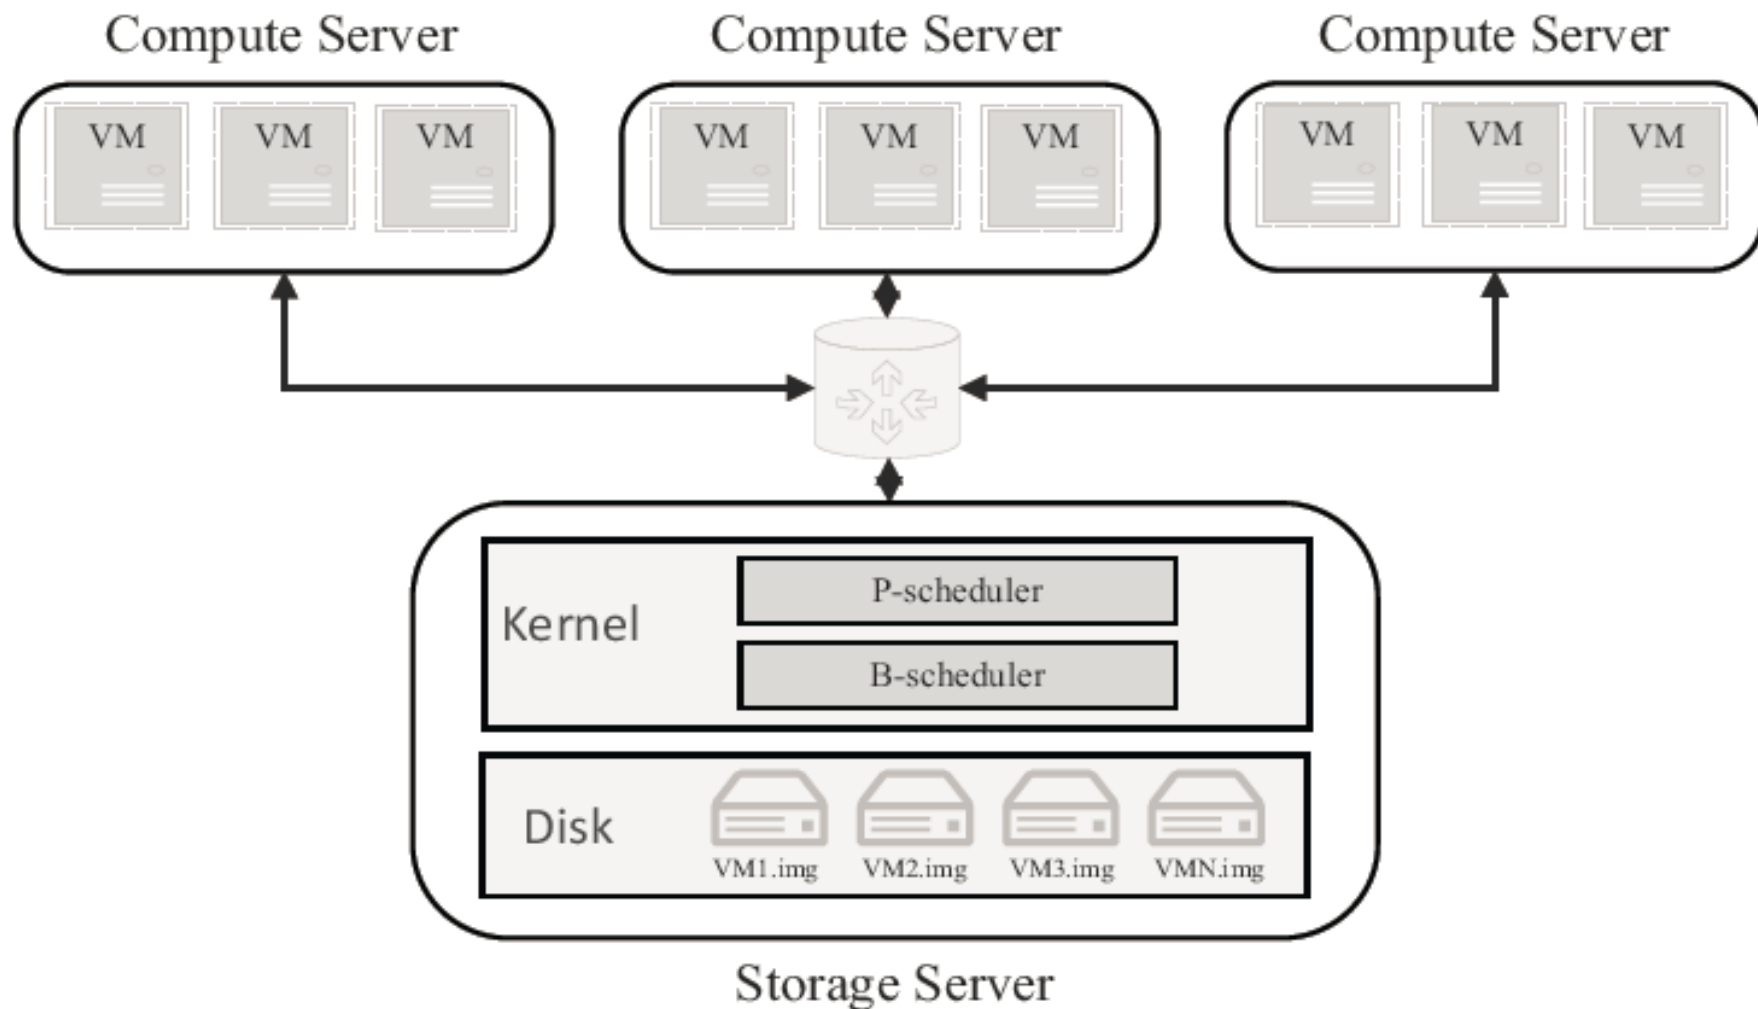


(a) P-scheduler



(b) B-scheduler


Testbed Deployment




Testbed Deployment

Tenant	Workload	Read (%)	IO Sizes	Seq/rand
A (VM1~VM5)	Index	75%	64KB	rand
B (VM6~VM10)	Data	61%	8KB	rand
C (VM11~VM15)	Message	56%	64KB	rand


SDFS: Evaluation (one-hypervisor)




Policy Enforcement with Multi-hypervisor




Policy Enforcement with Pulsar




Policy Enforcement with SDFS



SDFS vs. Pulsar



SDFS vs. Pulsar



Thanks!



lfzeng@hust.edu.cn