

数据中心技术

曾令仿、施展

武汉光电国家研究中心
华中科技大学
(2019-09-11 至 2019-11-01)



网络中心/计算中心/数据中心/?

High Performance Computing at the TU Kaiserslautern

Welcome to the homepage of the High Performance Computer **Elwetritsch** at the TU Kaiserslautern.

Aug. 29th: The number of abaqus and cae licenses has been enlarged.

The High Performance Computer 'Elwetritsch' combines the available HPC resources at the TU Kaiserslautern within the [Allianz für Hochleistungsrechnen Rheinland-Pfalz](#), abbr. **AHRP**. Access to the computer is granted to all members of the [TU Kaiserslautern](#) and the [AHRP](#).

Current Cluster Utilization

12670

38465

4198

Running

12670 job slots
1976 suspendable
(336 jobs from 46 users)

Pending

38465 job slots
746 suspendable
(1125 jobs from 13 users)

Available

4198 job slots on 185 hosts

Login Nodes	Users	Load /5min	Load /15min
TUK 1	69	0.99	0.84
TUK 2	66	1.14	0.80
TUK 3	24	2.20	2.22
TUK 4	81	2.80	2.74
RARP 1	2	0.10	0.13
RARP 2	1	0.08	0.12



Slots available: 16624 on 800 hosts



Hosts closed: 26 of 826 hosts



RAM: 71.2 TiB



/home: 9.6 TiB



/scratch: 1284.2 TiB

基本信息

➤ 课程主页

- <https://github.com/cs-course/data-center-course>

➤ 参考书

- 云计算与分布式系统——从并行处理到物联网，机械工业出版社，2012
- 云计算——概念、技术与架构，机械工业出版社，2014
- Barroso, Clidaras, and Holzle, “The Datacenter as a Computer: An Introduction to the Design of Warehouse-Scale Machines, Second Edition.”, 2013

基本信息

➤ 曾令仿

- 武汉光电国家研究中心，F302
- 课程公务 每周四上午08:30-10:00

➤ Email: lfzeng@hust.edu.cn

➤ 电话：027-87792450

➤ 个人主页：

<https://lingfangzeng.github.io/>

http://faculty.hust.edu.cn/zenglingfang/zh_CN/index.htm

基本信息

➤ 施展

- 武汉光电国家研究中心，F309
- 课程公务 每周一上午08:30-10:00

➤ Email: zshi@hust.edu.cn

➤ 电话：027-87792450

➤ 个人主页：

http://faculty.hust.edu.cn/shizhan/zh_CN/index.htm

授课目标

➤ 工程实践方面

- 能初步完成数据中心实际部署
 - OpenStack, Swarm, Mesos, Kubernetes (k8s), Lustre
- 具备运行、维护、使用基础技能
 - Linux, Bash, Object Storage, Vagrant, VirtualBox, Docker, Slurm

➤ 学术探索方面

- 熟悉相关领域前沿技术与进展
 - 新型半导体器件、分布式存储、虚拟化、软件定义数据中心
- 能独立开展相关领域创新性研究
 - 监控管理、调度迁移、应用、可靠、节能、安全

➤ 评分

- 分组研讨50%
- 实验报告50%

课程计划

序号	主题	日期
1	课程总体介绍、数据中心基础设施	9月11日
2	虚拟化、容器技术	9月18日
3	虚拟化环境存储优化	9月20日
4	软件定义数据中心	9月25日
5	新型非易失性存储器重塑数据中心	9月27日
6	大规模高性能分布式块存储系统数据中心部署实例	10月9日
7	监控与管理技术	10月11日
8	专题研讨	10月16日
9	专题研讨	10月18日
10	专题研讨	10月23日
11	专题研讨	10月25日
12	专题研讨	10月30日
13	专题研讨	11月1日

课程实践

➤ 形式与内容

- 围绕讲座与论文研讨内容选题
- 鼓励结合专业方向、兴趣特长自行设计，酌情优评
- 综合论文研讨与实验，形成报告

➤ 方法与环境

- 实验环境可以基于虚拟机、服务器
- 范例论文、基础实验方法课堂讲解
- 选题学习、实验重现课间课后分组进行

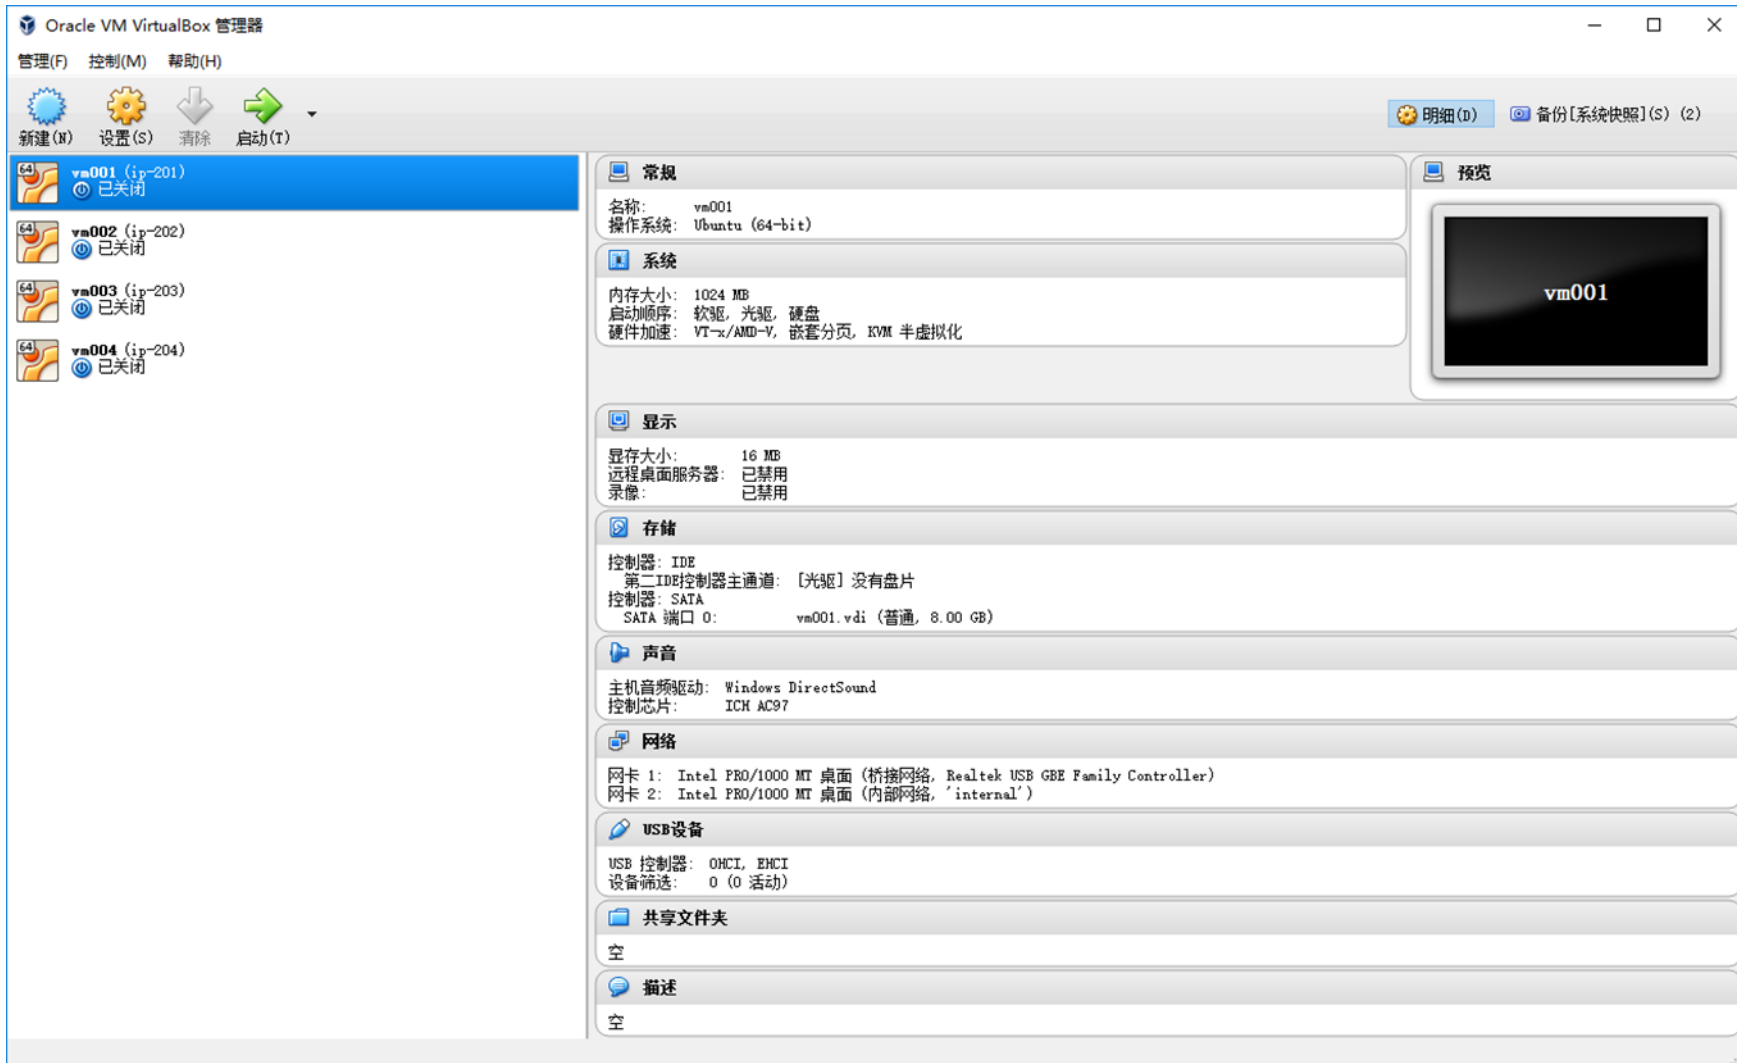
实验环境：物理机集群

- SSH远程访问
- 通过HTTP代理
- Ubuntu16.04LTS



机架服务器群 (国家光电研究中心F312)

实验环境：虚拟机集群



The screenshot displays the Oracle VM VirtualBox Manager interface. On the left, a list of virtual machines is shown, with 'vm001 (ip-201)' selected. The main pane on the right shows the configuration for 'vm001', which is currently powered off. The configuration is organized into several sections:

- 常规 (General):** Name: vm001, Operating System: Ubuntu (64-bit).
- 系统 (System):** Memory Size: 1024 MB, Boot Order: Floppy, CD-ROM, Hard Disk, Hardware Acceleration: VT-x/AMD-V, Nested Paging, KVM 半虚拟化.
- 显示 (Display):** Video Memory: 16 MB, Remote Desktop: Disabled, Recording: Disabled.
- 存储 (Storage):** Controller: IDE, Second IDE Controller Master: [Floppy] No Disk, Controller: SATA, SATA Port 0: vm001.vdi (Normal, 8.00 GB).
- 声音 (Audio):** Host Audio Driver: Windows DirectSound, Audio Chipset: ICH AC97.
- 网络 (Network):** Network Card 1: Intel PRO/1000 MT Desktop (Bridged network, Realtek USB GBE Family Controller), Network Card 2: Intel PRO/1000 MT Desktop (Internal network, 'internal').
- USB设备 (USB Devices):** USB Controller: OHCI, EHCI, Device Filtering: 0 (0 active).
- 共享文件夹 (Shared Folders):** Empty.
- 描述 (Description):** Empty.

On the far right, a 'Preview' window shows a black screen with the text 'vm001'.

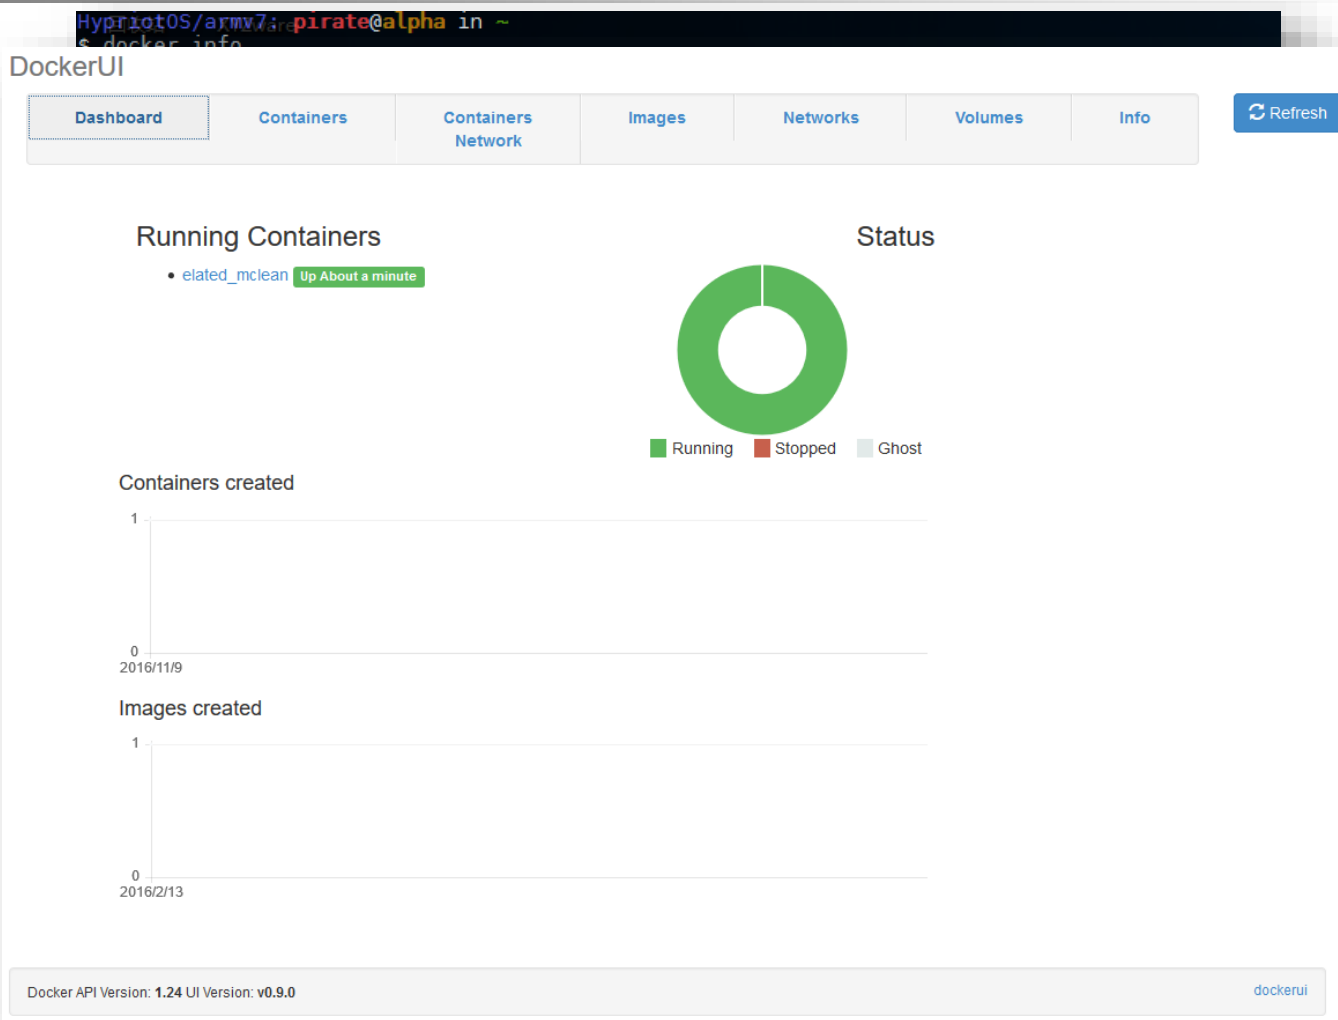
平台技术：虚拟机集群

```
Zhan@simba-thinkpad ~ $ ssh root@192.168.3.85
Welcome to Ubuntu 16.04.2 LTS (GNU/Linux 4.4.0-64-generic x86_64)
```

```
==> node3: Machine already provisioned. Run `vagrant provision` or use the `--provision` box)
==> node3: flag to force provisioning. Provisioners marked to run always will still run. (box)
==> node4: Clearing any previously set forwarded ports... poweroff (virtualbox)
==> node4: Fixed port collision for 22 => 2222. Now on port 2202.
==> node4: Clearing any previously set network interfaces...
==> node4: Preparing network interfaces based on configuration...
node4: Adapter 1: nat
node4: Adapter 2: hostonly
==> node4: Forwarding ports...
node4: 22 (guest) => 2202 (host) (adapter 1)
==> node4: Running 'pre-boot' VM customizations...
==> node4: Booting VM...
==> node4: Waiting for machine to boot. This may take a few minutes...
node4: SSH address: 127.0.0.1:2202
node4: SSH username: vagrant
node4: SSH auth method: private key
==> node4: Machine booted and ready!
==> node4: Checking for guest additions in VM...
==> node4: Setting hostname...
==> node4: Configuring and enabling network interfaces...
==> node4: Mounting shared folders...
node4: /vagrant => /root/vm-experiment/xenial-docker-cluster
node4: /vagrant_data => /root/vm-experiment/xenial-docker-cluster/data
==> node4: Machine already provisioned. Run `vagrant provision` or use the `--provision`
==> node4: flag to force provisioning. Provisioners marked to run always will still run.
controller xenial-docker-cluster #
```

```
==> node1: Clearing any previously set forwarded ports...
==> node1: Clearing any previously set network interfaces...
==> node1: Preparing network interfaces based on configuration...
```

平台技术：容器



课程QQ讨论群

- Σ 起来
- 扫码进群投票



群名称:数据中心技术2019

群 号:700415190